

Practical 8bis: Working with Text (Part 2)

The basics of Text Mining and NLP

Table of contents

1	Preamble	2
2	Setup	2
3	Illustrative Text Cleaning	6
4	Applying Normalisation	11
5	Process the Selected Listings	13
6	Frequencies and Ngrams	15
7	Count Vectoriser	16
8	TF/IDF Vectoriser	21
9	Word Clouds	22
10	Latent Dirichlet Allocation	24
11	Word2Vec	27
12	Processing the Full File	29

Part 2 of Practical 8 is *optional* and should only be attempted if Part 1 made sense to you.

1. The first few tasks are about finding important vocabulary (think ‘keywords’ and ‘significant terms’) in documents so that you can start to think about what is *distinctive* about documents and groups of documents. **This is quite useful and relatively easier to understand than what comes next!**
2. The second part is about fully-fledged NLP using Latent Dirichlet Allocation (topic modelling) and Word2Vec (words embeddings for use in clustering or similarity work).

The later parts are largely complete and ready to run; however, that *doesn't* mean you should just skip over them and think you've grasped what's happening and it will be easy to apply in your own analyses. I would *not* pay as much attention to LDA topic mining since I don't think its results are that good, but I've included it here as it's still commonly-used in the Digital Humanities and by Marketing folks. Word2Vec is much more powerful and forms the basis of the kinds of advances seen in ChatGPT and other LLMs.

Connections

Working with text is unquestionably *hard*. In fact, *conceptually* this is probably the most challenging practical of the term! But data scientists are *always* dealing with text because so much of the data that we collect (even more so thanks to the web) is not only text-based (URLs are text!) but, increasingly, unstructured (social media posts, tags, etc.). So while getting to grips with text is a challenge, it also uniquely positions you with respect to the skills and knowledge that other graduates are offering to employers.

1 Preamble

Parts of this practical have been written using `nltk` and `gensim`—both of which I have **removed** from the `sds2026` image to keep the size down and because relatively few students used any of this — but it would be *relatively* easy to rework it to use `spacy`. From what I can see, most programmers tend to use one *or* the other, and the switch wouldn't be hard, other than having to first load the requisite language models. You can [read about the models](#), and note that they are also [available in other languages](#) besides English.

2 Setup

Difficulty Level: Low

But this is only because this has been worked out for you. Starting from `sctach` in NLP is *hard* so people try to avoid it as much as possible.

2.1 Required Modules

Note

Notice that the number of modules and functions that we import is steadily increasing week-on-week, and that for text processing we tend to draw on quite a wide range of utilities! That said, the three most commonly used are: `sklearn` and `nltk`.

Standard libraries we've seen before.

```
import pandas as pd
import geopandas as gpd
import re
import matplotlib.pyplot as plt

# New, but fairly straightforward (for simple things)
from bs4 import BeautifulSoup
```

Vectorisers we will use from the 'big beast' of Python machine learning: Sci-Kit Learn.

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.decomposition import LatentDirichletAllocation

# We don't use this but I point out where you *could*
from sklearn.preprocessing import OneHotEncoder

```

NLP-specific libraries that we will use for tokenisation, lemmatisation, and frequency analysis.

```

import nltk

try:
    from nltk.corpus import wordnet as wn
    from nltk.stem.wordnet import WordNetLemmatizer
    from nltk.corpus import stopwords
    stopwords.words('english')
    lemmatizer = WordNetLemmatizer()
    tokenizer = ToktokTokenizer()
except:
    nltk.download("stopwords")
    nltk.download("punkt_tab")
    nltk.download("averaged_perceptron_tagger_eng")
    from nltk.corpus import stopwords
    from nltk.corpus import wordnet as wn
    from nltk.stem.wordnet import WordNetLemmatizer
    stopwords.words('english')

from nltk.corpus import stopwords
stopword_list = set(stopwords.words('english'))

from nltk.tokenize import word_tokenize, sent_tokenize
from nltk.tokenize.toktok import ToktokTokenizer

from nltk.stem.porter import PorterStemmer
from nltk.stem.snowball import SnowballStemmer

from nltk import ngrams, FreqDist

lemmatizer = WordNetLemmatizer()
tokenizer = ToktokTokenizer()

```

```

[nltk_data] Downloading package stopwords to
[nltk_data]      /Users/jreades/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt_tab to
[nltk_data]      /Users/jreades/nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger_eng to
[nltk_data]      /Users/jreades/nltk_data...
[nltk_data] Package averaged_perceptron_tagger_eng is already up-to-
[nltk_data]      date!

```

Remaining libraries that we'll use for processing and display text data. Most of this relates to dealing with the various ways that text data cleaning is *hard* because of the myriad formats it comes in.

```
try:
    from wordcloud import WordCloud
except:
    ! pip install wordcloud
import string
import unicodedata
from wordcloud import WordCloud, STOPWORDS
```

This next is just a small utility function that allows us to output Markdown (like this cell) instead of plain text:

```
from IPython.display import display_markdown

def as_markdown(head='', body='Some body text'):
    if head != '':
        display_markdown(f"##### {head}\n\n{body}\n", raw=True)
    else:
        display_markdown(f">{body}\n", raw=True)

as_markdown('Result!', "Here's my output...")
```

Result!

Here's my output...

2.2 Loading Data

Connections

Because I generally want each practical to stand on its own (unless I'm trying to make a *point*), I've not moved this to a separate Python file (e.g. `utils.py`, but in line with what we covered back in the lectures on [Functions and Packages](#), this sort of thing is a good candidate for being split out to a separate file to simplify re-use.

Remember this function from last week? We use it to save downloading files that we already have stored locally. But notice I've made some small changes... what do these do to help the user?

```
from pathlib import Path
from requests import get
from functools import wraps

def check_cache(f):
    @wraps(f)
    def wrapper(src:str, dst:str, min_size=100) -> Path:
```

```

    if src.find('?') == -1:
        url = Path(src)
    else:
        url = Path(src[:src.find('?')])
    fn = url.name # Extract the filename
    dsn = Path(f"{dst}/{fn}") # Destination filename
    if dsn.is_file() and dsn.stat().st_size > min_size:
        print(f"+ {dsn} found locally!")
        return(dsn)
    else:
        print(f"+ {dsn} not found, downloading!")
        return(f(src, dsn))
return wrapper

```

@check_cache

```

def cache_data(src:Path, dst:Path) -> str:
    """Downloads a remote file.

```

The function sits between the 'read' step of a pandas or geopandas data frame and downloading the file from a remote location. The idea is that it will save it locally so that you don't need to remember to do so yourself. Subsequent re-reads of the file will return instantly rather than downloading the entire file for a second or n-th itme.

Parameters

src : str

The remote **source** for the file, any valid URL should work.

dst : str

The **destination** location to save the downloaded file.

Returns

str

A string representing the local location of the file.

"""

```

# Create any missing directories in dest(ination) path
# -- os.path.join is the reverse of split (as you saw above)
# but it doesn't work with lists... so I had to google how
# to use the 'splat' operator! os.makedirs creates missing
# directories in a path automatically.

```

```

if not dst.parent.exists():
    dst.parent.mkdir(parents=True, exist_ok=True)

```

```

# Download and write the file

```

```

with dst.open(mode='wb') as file:
    response = get(src)
    file.write(response.content)

```

```
print(' + Done downloading...')

return dst.resolve()
```

Tip

For very large *non-geographic* data sets, remember that you can use `use_cols` (or `columns` depending on the file type) to specify a subset of columns to load.

Load the main data set:

```
# Load the data sets created in the previous practical
lux    = gpd.read_parquet(Path('data/clean/luxury.geoparquet'))
aff    = gpd.read_parquet(Path('data/clean/affordable.geoparquet'))
bluesp = gpd.read_parquet(Path('data/clean/bluespace.geoparquet'))
```

3 Illustrative Text Cleaning

Now we're going to step through the *parts* of the process that we apply to clean and transform text. We'll do this individually before using a function to apply them *all at once*. To keep things moving along, we're going to take a small sample from the data set.

```
pd.set_option("display.max_colwidth", 250)
sample = bluesp.description.sample(5, random_state=44)
sample
```

```
10390    Stunning & Spacious Air Conditioned 3 bedroom 3BR Apartment on the Ground floor
34253    his charming house boasts a prime location just a 5-minute stroll from iconic
49617    Key features:<br />Three double bedrooms<br />Two bathrooms<br />Close to Water
56882    Private double ensuite bedroom in a shared home. 8 min walk to Surbiton station
21289    Cosy two bedroom apartment on the ninth floor with an amazing river view, just
Name: description, dtype: object
```

3.1 Removing HTML

Difficulty level: Moderate

You'll have seen that there are some things that are *not* text in the sample, most notably this stuff: `<br>` and `<br />`. That is raw HTML, the markup language for web pages. To go any further we will need to strip this out. Depending on what you needed to extract from a web page you *can* select different elements by name, class, or id, and do a whole bunch of other things, but here we're just going to crudely strip the tags out.

You're also going to see a powerful new method: `apply`. `Apply` simply runs some function against every row or column in the data set depending on how it's called. In this case, we use a `lambda` function (a function without a formal definition) to create a 'Beautiful Soup' HTML parser and then strip the tags from each row.

Hint: you need to need to **get the text** out of the each returned `<p>` and `<div>` element! I'd suggest also commenting this up since there is a *lot* going on on some of these lines of code!

Question

```
cleaned = sample.apply(lambda x: BeautifulSoup(??, 'html.parser').??(" "))
cleaned
```

```
10390    Stunning & Spacious Air Conditioned 3 bedroom 3BR Apartment on the Ground floor
34253    his charming house boasts a prime location just a 5-minute stroll from iconic
49617    Key features: Three double bedrooms Two bathrooms Close to Waterloo Excellent
56882    Private double ensuite bedroom in a shared home. 8 min walk to Surbiton station
21289    Cosy two bedroom apartment on the ninth floor with an amazing river view
Name: description, dtype: object
```

3.2 Lower Case

 Difficulty Level: Low.

Question

```
lower = cleaned.??
lower
```

You should get something like the following:

```
10390    stunning & spacious air conditioned 3 bedroom 3br apartment on the ground floor
34253    his charming house boasts a prime location just a 5-minute stroll from iconic
49617    key features: three double bedrooms two bathrooms close to waterloo excellent
56882    private double ensuite bedroom in a shared home. 8 min walk to surbiton station
21289    cosy two bedroom apartment on the ninth floor with an amazing river view
Name: description, dtype: object
```

3.3 Stripping 'Punctuation'

 Difficulty level: Hard

This is because you need to understand: 1) why we're *compiling* the regular expression and how to use character classes; and 2) how the NLTK tokenizer differs in its approach to the regex.

3.3.1 Regular Expression Approach

We want to clear out punctuation using a regex that takes advantage of the `[...]` (character class) syntax. The really tricky part is remembering how to specify the ‘punctuation’ when some of that punctuation has ‘special’ meanings in a regular expression context. For instance, `.` means ‘any character’, while `[` and `]` mean ‘character class’. So this is another *escaping* problem and it works the *same* way it did when we were dealing with the Terminal...

Hints: some other factors...

1. You will want to match more than one piece of punctuation at a time, so I’d suggest add a `+` to your pattern.
2. You will need to look into *metacharacters* for creating a kind of ‘any of the characters *in this class*’ bag of possible matches.

Question

```
pat = re.compile(r'[???]+')
subbed = lower.apply(lambda x: pat.sub('',x))
```

Depending on how thorough you are, you should get something like this:

```
10390    stunning    spacious air conditioned 3 bedroom 3br apartment on the ground floor
34253    his charming house boasts a prime location just a 5minute stroll from iconic l
49617    key features: three double bedrooms two bathrooms close to waterloo excellent
56882    private double ensuite bedroom in a shared home 8 min walk to surbiton station
21289                                cosy two bedroom apartment on the ninth floor with an amazing r
Name: description, dtype: object
```

3.3.2 Tokenizer

The other way to do this, which is probably *easier*, is to draw on the tokenizers [already provided by NLTK](#). For our purposes `word_tokenize` is probably fine, but depending on your needs there are other options and you can also write your own.

```
nltk.download('punkt')
nltk.download('wordnet')
from nltk.tokenize import word_tokenize
tokens = lower.apply(word_tokenize)
tokens
```

```
[nltk_data] Downloading package punkt to /Users/jreades/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /Users/jreades/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

```
10390    [stunning, &, spacious, air, conditioned, 3, bedroom, 3br, apartment, on, the,
34253    [his, charming, house, boasts, a, prime, location, just, a, 5-
minute, stroll, from, iconic, landmarks, such, as, big, ben, and, the, london, eye, ,,
49617    [key, features, :, three, double, bedrooms, two, bathrooms, close, to, waterloo
```

```
56882      [private, double, ensuite, bedroom, in, a, shared, home, ., 8, min, walk, to,
21289      [cosy, two, bedroom, apartment, on, the, ninth, floor, with, an, amazing, rive
Name: description, dtype: object
```

⚠ Not Equivalent!

Notice that these aren't giving you the same result!

For our purposes it is easier in the subsequent stages to work with the `list` output from `tokenize` rather than the string output from `re.sub`, but if you needed to convert tokenise back to a string you could do it this way:

```
tokens = lower.apply(word_tokenize).apply(lambda x: ' '.join(x))
```

The other way around would be:

```
lower.apply(lambda x: pat.sub(',',x)).str.split(' ')
```

3.4 Stopword Removal

⚠ Difficulty Level: Moderate

You need to remember how list comprehensions work to use the `stopword_list`.

```
stopword_list = set(stopwords.words('english'))
print(stopword_list)
```

```
{'or', 'how', 'they', 'in', 'whom', 'down', 'by', 'its', "couldn't", "they'd", 'any', ...}
```

Question

```
passthru = []
for t in tokens[2:4]:
    passthru.append([x for x in t if ?? and len(x) > 1])

for p in passthru:
    as_markdown("Stopwords removed:", p)
```

3.5 Lemmatisation vs Stemming

💡 Difficulty level: Low.

```
from nltk.stem.porter import PorterStemmer
from nltk.stem.snowball import SnowballStemmer
from nltk.stem.wordnet import WordNetLemmatizer
```

```

lemmatizer = WordNetLemmatizer()
print(lemmatizer.lemmatize('monkeys'))
print(lemmatizer.lemmatize('cities'))
print(lemmatizer.lemmatize('complexity'))
print(lemmatizer.lemmatize('Reades'))

```

monkey
 city
 complexity
 Reades

```

stemmer = SnowballStemmer(language='english')
print(stemmer.stem('monkeys'))
print(stemmer.stem('cities'))
print(stemmer.stem('complexity'))
print(stemmer.stem('Reades'))

```

monkey
 citi
 complex
 read

```

lemmatizer = WordNetLemmatizer()
stemmer = SnowballStemmer(language='english')

# This would be better if we passed in a PoS (Part of Speech) tag as well,
# but processing text for parts of speech is *expensive* and for the purposes
# of this tutorial, not necessary.
lemmas = passthru.apply(lambda x: [lemmatizer.lemmatize(t) for t in x])
stemmed = passthru.apply(lambda x: [stemmer.stem(t) for t in x])

as_markdown("Lemmatised", '\n\n>'.join(lemmas.str.join(' ').to_list()))
print("\n\n")
as_markdown("Stemmed", '\n\n>'.join(stemmed.str.join(' ').to_list()))

```

Lemmatised

stunning spacious air conditioned bedroom 3br apartment ground floor terrace prime
 location chelsea perfect place lovely comfortable stay heart chelsea minute walk
 river thames stamford bridge short journey bus underground south kensington earl
 court knightsbridge hyde park easy access london attraction perfect choice family
 group friend

charming house boast prime location 5-minute stroll iconic landmark big ben london
 eye well conveniently close waterloo station 'll find bus stop supermarket nearby
 added convenience inside house comprises three cozy bedroom one luxurious queen-
 sized bed two two comfortable single bed kitchen equipped basic cooking one toilet
 sink well full bathroom

key feature three double bedroom two bathroom close waterloo excellent transport link beautiful finish wood flooring throughout full description impressive 1115sqft 104sqm contemporary three double bedroom two bathroom first floor conversion kennington road short walk waterloo lambeth north station cut old vic 's bar restaurant

private double ensuite bedroom shared home min walk surbiton station 20 min direct waterloo min walk river thames 20 min walk kingston town centre easy bus train wimbledon home shared least two people time sharing kitchen space room bathroom attached room house listed separately 1-2 guest time able see previous review listing one new

cosy two bedroom apartment ninth floor amazing river view minute pimlico station heart city zone bigben buckingham palace hyde park chelsea 15/20 minute walking distance

Stemmed

stun spacious air condit bedroom 3br apart ground floor terrac prime locat chelsea perfect place love comfort stay heart chelsea minut walk river thame stamford bridg short journey bus underground south kensington earl court knightsbridg hyde park easi access london attract perfect choic famili group friend

charm hous boast prime locat 5-minut stroll icon landmark big ben london eye well conveni close waterloo station ll find bus stop supermarket nearbi ad conveni insid hous compris three cozi bedroom one luxuri queen-siz bed two two comfort singl bed kitchen equip basic cook one toilet sink well full bathroom

key featur three doubl bedroom two bathroom close waterloo excel transport link beauti finish wood floor throughout full descript impress 1115sqft 104sqm contemporari three doubl bedroom two bathroom first floor convers kennington road short walk waterloo lambeth north station cut old vic 's bar restaur

privat doubl ensuit bedroom share home min walk surbiton station 20 min direct waterloo min walk river thame 20 min walk kingston town centr easi bus train wimbledon home share least two peopl time share kitchen space room bathroom attach room hous list separ 1-2 guest time abl see previous review list one new

cosi two bedroom apart ninth floor amaz river view minut pimlico station heart citi zone bigben buckingham palac hyde park chelsea 15/20 minut walk distanc

4 Applying Normalisation

The above approach is fairly hard going since you need to loop through every list element applying these changes one at a time. Instead, we could convert the column to a corpus (or use `pandas apply`) together with a function imported from a library to do the work.

4.1 Downloading the Custom Module

In a Jupyter notebook, this code allows us to edit and reload the library dynamically:

```
%load_ext autoreload
%autoreload 2
```

💡 Difficulty level: Low.

This custom module is not perfect, but it gets the job done... mostly and has some additional features that you could play around with for a final project (e.g. `detect_entities` and `detect_acronyms`).

```
try:
    from textual import *
except:
    try:
        from unicode import unicode
    except:
        ! pip install unicode
    import urllib.request
    host = 'http://orca.casa.ucl.ac.uk'
    turl = f'{host}/~jreades/__textual__.py'
    tdirs = Path('textual')
    tpath = Path(tdirs / '__init__.py')

    if not tpath.exists():
        tdirs.mkdir(parents=True, exist_ok=True)
        urllib.request.urlretrieve(turl, tpath)
    from textual import *
```

All NLTK libraries installed...

4.2 Importing the Custom Module

💡 Difficulty Level: Low.

But only because you didn't have to write the module! However, the questions could be hard...

```
normalised = sample.apply(normalise_document, remove_digits=True)
as_markdown('Normalised', '\n\n>'.join(normalised.to_list()))
```

Normalised

stunning spacious conditioned bedroom apartment ground floor terrace prime location chelsea . perfect place lovely comfortable stay heart chelsea minute walk river

thames stamford bridge short journey underground south kensington earls court
knightsbridge hyde park easy access london attractions . perfect choice families
group friends

charm house boast prime location minute stroll iconic landmark london well conve-
niently close waterloo station . find stop supermarket nearby added convenience
. inside house comprise three cozy bedrooms luxurious queensized comfortable
single . kitchen equip basic cooking toilet sink well full bathroom .

feature three double bedroom bathroom close waterloo excellent transport link
beautiful finish wood flooring throughout full description impressive sqft . contem-
porary three double bedroom bathroom first floor conversion kennington road short
walk waterloo lambeth north station restaurant .

private double ensuite bedroom share home . walk surbiton station . direct waterloo
. walk river thames walk kingston town centre easy train wimbledon . home share
least people time share kitchen space . room bathroom attach . room house list
separately guest time . able previous review listing

cosy bedroom apartment ninth floor amazing river view minute pimlico station heart
city . zone . bigben buckingham palace hyde park chelsea minute walk distance .

```
help(normalise_document)
```

🔥 Stop!

Beyond this point, we are moving into Natural Language Processing. If you are already struggling with regular expressions, I would recommend *stopping here*. You can come back to revisit the NLP components and creation of word clouds later.

5 Process the Selected Listings

💡 Difficulty level: Low, but you'll need to be patient!

Notice the use of `%%time` here – this will tell you how long each block of code takes to complete. It's a really useful technique for reminding *yourself* and others of how long something might take to run. I find that with NLP this is particularly important since you have to do a *lot* of processing on each document in order to normalise it.

💡 Tip

Notice how we can change the default parameters for `normalise_document` even when using `apply`, but that the syntax is different. So whereas we'd use `normalise_document(doc, remove_digits=True)` if calling the function directly, here it's `.apply(normalise_document, remove_digits=True)`!

Question

```
%%time
# I get about 2 seconds on a M2 Mac
bluesp['description_norm'] = bluesp.???apply(???, remove_digits=True)
```

5.1 Select and Tokenise

💡 Difficulty level: Low, except for the double list-comprehension.

5.1.1 Select and Extract Corpus

See useful tutorial [here](#). Although we shouldn't have any empty descriptions, by the time we've finished normalising the textual data we may have *created* some empty values and we need to ensure that we don't accidentally pass a NaN to the vectorisers and frequency distribution functions.

```
# This makes it easy to go back and re-run
# the analysis with other data subsets.
srcdf = bluesp
```

💡 Coding Tip

Notice how you only need to change the value of the variable here to try any of the different selections we did above? This is a simple kind of parameterisation somewhere between a function and hard-coding everything.

```
corpus = srcdf.description_norm.fillna(' ').values
print(corpus[0:3])
```

```
['well decorate sunny garden flat fulham park gardens close putney bridge tube .this co
' enjoy warm cosy modern apartment . river view gorgeous view city . walk barking stati
' make memory unique familyfriendly place . near nice bridge minute away harrods london
```

5.1.2 Tokenise

There are different forms of tokenisation and different algorithms will expect differing inputs. Here are two:

```
sentences = [nltk.sent_tokenize(text) for text in corpus]
words      = [[nltk.tokenize.word_tokenize(sentence)
               for sentence in nltk.sent_tokenize(text)]
              for text in corpus]
```

Notice how this has turned every sentence into an array and each document into an array of arrays:

```
print(f"Sentences 0: {sentences[0]}")
print()
print(f"Words 0: {words[0]}")
```

Sentences 0: ['well decorate sunny garden flat fulham park gardens close putney bridge

Words 0: [['well', 'decorate', 'sunny', 'garden', 'flat', 'fulham', 'park', 'gardens',

6 Frequencies and Ngrams

 Difficulty level: Moderate.

One new thing you'll see here is the `ngram`: ngrams are 'simply' pairs, or triplets, or quadruplets of words. You may come across the terms `unigram` (`ngram(1,1)`), `bigram` (`ngram(2,2)`), `trigram` (`ngram(3,3)`)... typically, you will rarely find anything beyond trigrams, and these present real issues for `text2vec` algorithms because the embedding for geographical information, and systems is *not* the same as for geographical information systems.

6.0.1 Build Frequency Distribution

Build counts for ngram range 1..3:

```
fcounts = dict()

# Here we replace all full-stops... can you think why we might do this?
data = nltk.tokenize.word_tokenize(' '.join([text.replace('.', '') for text in corpus]))

for size in 1, 2, 3:
    fdist = FreqDist(ngrams(data, size))
    print(fdist)
    # If you only need one note this: https://stackoverflow.com/a/52193485/4041902
    fcounts[size] = pd.DataFrame.from_dict({'Ngram Size {size}': fdist})
```

<FreqDist with 1838 samples and 13751 outcomes>

<FreqDist with 8020 samples and 13750 outcomes>

<FreqDist with 10515 samples and 13749 outcomes>

6.0.2 Output Top-n Ngrams

And output the most common ones for each ngram range:

```
for dfs in fcounts.values():
    print(dfs.sort_values(by=dfs.columns.values[0], ascending=False).head(10))
    print()
```

Ngram Size 1	
walk	501
london	327
minute	282
river	276
station	275
view	194
apartment	193
thames	172
room	164
bedroom	162

Ngram Size 2		
minute	walk	190
river	view	98
	thames	83
central	london	74
canary	wharf	63
walk	river	46
	waterloo	42
thames	river	41
living	room	37
walk	distance	36

Ngram Size 3			
walk	river	thames	37
station	minute	walk	25
fully	equip	kitchen	23
walk	waterloo	station	20
close	river	thames	18
within	walk	distance	17
minute	walk	river	17
walk	thames	river	17
minute	walk	waterloo	14
thames	river	view	13

Questions

1. Can you think why we don't care about punctuation for frequency distributions and n-grams?
2. Do you understand what n-grams *are*?

7 Count Vectoriser

💡 Difficulty level: Low, but the output needs some thought!

This is a big foray into sklearn (sci-kit learn) which is the main machine learning and clustering module for Python. For processing text we use a *vectorisers* to convert terms to a vector

representation. We're doing this on the smallest of the derived data sets because these processes can take a while to run and generate *huge* matrices (remember: one row and one column for each term!).

7.1 Fit the Vectoriser

```
cvectorizer = CountVectorizer(ngram_range=(1,3))
cvectorizer.fit(corpus)
```

```
CountVectorizer(ngram_range=(1, 3))
```

7.2 Brief Demonstration

Once we've fit the vectoriser, we can find the integer associated with any word in the corpus and can also ask how many times it occurs:

```
pd.options.display.max_colwidth=750

# Note that the term can be a unigram, a bigram, or a trigram!
term = 'minute walk'
# Find the vocabulary mapping for the term
print(f"Vocabulary mapping for '{term}' is {cvectorizer.vocabulary_[term]}")

# How many times is it in the data
print(f"Found {srcdf.description_norm.str.contains(term).sum():,} rows containing '{term}'")

# Print the descriptions containing the term
for i, x in enumerate(srcdf[srcdf.description_norm.str.contains(term)].description_norm):
    as_markdown(f"{term}: {i}", x)
    if i > 4: break
```

```
Vocabulary mapping for 'minute walk' is 10349
Found 115 rows containing 'minute walk'
```

minute walk: 0

make memory unique familyfriendly place . near nice bridge minute away harrods london . cycling area jungle park river view playground available walkable distance . minute walk uxbridge underground station . minute walk shop mall . private parking bike . security service available . security deposit pound . refund hour checkout .

minute walk: 1

nice cozy private room pimlico central london . prime location visit city close river thames minute walk pimlico station walk victoria station . stop local shop within minute walk . apartment basic clean proper picture apartment soon .

minute walk: 2

private quiet bright large double room ensuite bathroom . location stun next thames
minute walk stamford brook tube . close riverside great restaurant place historic
interest .

minute walk: 3

beautiful modern onebedroom flat heart london . covent garden literally minute
walk apartment . charring cross embankment covent garden underground station
minute walk . plenty restaurant around minute walk thames river . london .

minute walk: 4

people stay love place quirt private cosy serene culture easy transport link whole
london . bedroom double brand suit . hammock relax front room . minute walk
canal cafe restaurant near local crown meter walk flat

minute walk: 5

suits couple front room maybe minute elephant castle tube bath shower minikitchen
small garden small dining room . minute walk thames minute london bridge

7.3 Transform the Corpus

You can only *transform* the entire corpus *after* the vectoriser has been fitted. There is an option to `fit_transform` in one go, but I wanted to demonstrate a few things here and some vectorisers are don't support the one-shot fit-and-transform approach. **Note the type of the transformed corpus:**

```
cvtcorpus = cvvectorizer.transform(corpus)
cvtcorpus # cvtcorpus for count-vectorised transformed corpus
```

```
<347x19391 sparse matrix of type '<class 'numpy.int64'>'
  with 38234 stored elements in Compressed Sparse Row format>
```

7.3.1 Single Document

Here is the **first** document from the corpus:

```
doc_df = pd.DataFrame(cvtcorpus[0].T.todense(),
                      index=cvvectorizer.get_feature_names_out(), columns=["Counts"])
                      ).sort_values('Counts', ascending=False)
doc_df.head(10)
```

	Counts
close	3

	Counts
flat	2
park	2
area	2
tennis	1
park tennis	1
comfortable	1
gardens	1
shop	1
residential street	1

7.3.2 Transformed Corpus

```
cvdf = pd.DataFrame(data=cvtcorpus.toarray(),
                    columns=cvectorizer.get_feature_names_out())
print(f"Raw count vectorised data frame has {cvdf.shape[0]:,} rows and {cvdf.shape[1]:,} columns")
cvdf.iloc[0:3,0:7]
```

Raw count vectorised data frame has 347 rows and 19,391 columns.

	aback	aback spacious	aback spacious apartment	abba	abba voyage	abba voyage walk	abba wa
0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0

```
cvdf.iloc[-3:,-7:]
```

	zone public transportation	zone recently	zone recently refurbish	zone tube	zone tube station	z
344	0	0	0	0	0	0
345	0	0	0	0	0	0
346	0	0	0	0	0	0

Questions

1. Why is the [single document](#) view a list of terms and counts?
2. Why is the [corpus view](#) a table of terms (columns) and integers (rows)?

7.3.3 Filter Low-Frequency Words

These are likely to be artefacts of text-cleaning or human input error. As well, if we're trying to look across an entire corpus then we might not want to retain words that only appear in a couple of documents.

Let's start by getting the *column* sums:

```

sums = cvdf.sum(axis=0)
print(f"There are {len(sums):,} terms in the data set.")
sums.head()

```

There are 19,391 terms in the data set.

```

aback          1
aback spacious 1
aback spacious apartment 1
abba           2
abba voyage    1
dtype: int64

```

Remove columns (i.e. terms) appearing in less than 1% of documents. You can do this by thinking about what the shape of the data frame means (rows and/or columns) and how you'd get 1% of that!

Question

```

filter_terms = sums >= cvdf.shape[0] * ???

```

Now see how we can use this to strip out the columns corresponding to low-frequency terms:

```

fcvdf = cvdf.drop(columns=cvdf.columns[~filter_terms].values)
print(f"Filtered count vectorised data frame has {fcvdf.shape[0]:,} rows and {fcvdf.shape[1]:,} columns.")
fcvdf.iloc[0:3,0:7]

```

Filtered count vectorised data frame has 347 rows and 1,688 columns.

	access	access buses	access buses west	access central	access central london	access everything	access everything else
0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0

```

fcvdf.sum(axis=0)

```

```

access          68
access buses     4
access buses west 4
access central   5
access central london 5
..
zone comfortable cosy 7
zone london         6
zone london warm    4
zone recently       8
zone recently refurbish 8
Length: 1688, dtype: int64

```

We're going to pick this up again in Task 7.

Questions

- Can you explain what `doc_df` contains?
- What does `cvdf` contain? Explain the rows and columns.
- What is the function of `filter_terms`?

8 TF/IDF Vectoriser

 Difficulty level: Moderate

But only if you want to understand how `max_df` and `min_df` work!

8.1 Fit and Transform

```
tfvectorizer = TfidfVectorizer(use_idf=True, ngram_range=(1,3),
                               max_df=0.75, min_df=0.01) # <-- these matter!
tftcorpus    = tfvectorizer.fit_transform(corpus) # TF-transformed corpus
```

8.2 Single Document

```
doc_df = pd.DataFrame(tftcorpus[0].T.todense(), index=tfvectorizer.get_feature_names_out())
doc_df.sort_values('Weights', ascending=False).head(10)
```

	Weights
close	0.243962
area	0.207529
quiet residential	0.203155
putney	0.203155
street close	0.203155
kings road	0.196091
residential	0.184943
decorate	0.184943
tennis	0.184943
sunny	0.184943

8.3 Transformed Corpus

```
tfidf = pd.DataFrame(data=tftcorpus.toarray(),
                     columns=tfvectorizer.get_feature_names_out())
print(f"TF/IDF data frame has {tfidf.shape[0]:,} rows and {tfidf.shape[1]:,} columns")
tfidf.head()
```

TF/IDF data frame has 347 rows and 1,660 columns.

	access	access buses	access buses west	access central	access central london	access everything	a
0	0.0	0.0	0.0	0.0	0.0	0.0	0
1	0.0	0.0	0.0	0.0	0.0	0.0	0
2	0.0	0.0	0.0	0.0	0.0	0.0	0
3	0.0	0.0	0.0	0.0	0.0	0.0	0
4	0.0	0.0	0.0	0.0	0.0	0.0	0

Questions

- What does the TF/IDF score *represent*?
- What is the role of `max_df` and `min_df`?

9 Word Clouds

9.1 For Counts

💡 Difficulty level: Easy!

```
fcvdf.sum().sort_values(ascending=False)
```

```
walk          501
london        327
minute        282
river         276
station       275
...
community     4
common easy access 4
common easy   4
ready         4
retreat rest   4
Length: 1688, dtype: int64
```

```
ff = Path('RobotoSlab-VariableFont_wght.ttf')
dp = Path('/home/jovyan/.fonts/')
tp = Path(Path.home(), 'Library', 'Fonts')
if tp.exists():
    fp = tp / ff
else:
    fp = dp / ff
```

```
f,ax = plt.subplots(1,1,figsize=(5,5))
plt.gcf().set_dpi(300)
Cloud = WordCloud(
    background_color="white",
    max_words=75,
    font_path=fp
).generate_from_frequencies(fcvdf.sum())
ax.imshow(Cloud)
ax.axis("off");
#plt.savefig("Wordcloud 1.png")
```



9.2 For TF/IDF Weighting

💡 Difficulty level: Low, but you'll need to be patient!

```
tfidf.sum().sort_values(ascending=False)
```

```
walk                23.946637
minute              17.656652
london              17.114220
river               15.509969
apartment           14.143728
...
smart bedroom       0.332017
smart bedroom battersea 0.332017
smart netflixs close 0.332017
terrace overlooking 0.332017
netflixs close      0.332017
Length: 1660, dtype: float64
```

```
f,ax = plt.subplots(1,1,figsize=(5,5))
plt.gcf().set_dpi(300)
Cloud = WordCloud(
    background_color="white",
```

```
max_words=100,  
font_path=fp  
) .generate_from_frequencies(tfidf.sum())  
ax.imshow(Cloud)  
ax.axis("off");  
#plt.savefig("Wordcloud 2.png")
```



Questions

- What does the `sum` represent for the count vectoriser?
- What does the `sum` represent for the TF/IDF vectoriser?

10 Latent Dirichlet Allocation

 Tip

I would give this a *low* priority. It's a commonly-used method, but on small data sets it really isn't much use and I've found its answers to be... unclear... even on large data sets.

I would give this a *low* priority. It's a commonly-used method, but on small data sets it really isn't much use and I've found its answers to be... unclear... even on large data sets.

Adapted from [this post](#) on doing LDA using sklearn. Most other examples use the `gensim` library.

```
# Notice change to ngram range
# (try 1,1 and 1,2 for other options)
vectorizer = CountVectorizer(ngram_range=(1,2))
```

10.1 Calculate Topics

```
vectorizer.fit(corpus)
tcorpus = vectorizer.transform(corpus) # tcorpus for transformed corpus

LDA = LatentDirichletAllocation(n_components=3, random_state=42) # Might
LDA.fit(tcorpus)
```

```
LDA = LatentDirichletAllocation(n_components=3, random_state=42) # Might want to exp
LDA.fit(tcorpus)
```

```
LatentDirichletAllocation(n_components=3, random_state=42)
```

```
first_topic = LDA.components_[0]
top_words = first_topic.argsort()[-25:]

for i in top_words:
    print(vectorizer.get_feature_names_out()[i])
```

minute walk
easy
stay
locate
access
wharf
waterloo
park
canal
floor
tube
thames
river view
balcony
close
minute
apartment
bedroom
room
river
flat
view
station
london
walk

```
for i, topic in enumerate(LDA.components_):
    as_markdown(f'Top 10 words for topic #{i}', ', '.join([vectorizer.get_feature_names_out()[i] for i in top_words]))
```

Top 10 words for topic #0

minute walk, easy, stay, locate, access, wharf, waterloo, park, canal, floor, tube, thames, river view, balcony, close, minute, apartment, bedroom, room, river, flat, view, station, london, walk

Top 10 words for topic #1

kitchen, available, shop, stay, wharf, zone, room, away, waterloo, great, flat, thames, restaurant, view, location, city, close, central, river, minute walk, apartment, station, walk, london, minute

Top 10 words for topic #2

location, restaurant, central london, shop, double, river view, bathroom, kitchen, private, river thames, park, central, flat, close, apartment, room, bedroom, view, minute walk, thames, minute, station, london, river, walk

10.2 Maximum Likelihood Topic

```
topic_values = LDA.transform(tcorpus)
topic_values.shape
```

(347, 3)

```
pd.options.display.max_colwidth=20
srcdf['Topic'] = topic_values.argmax(axis=1)
srcdf.head()
```

	geometry	listing_url	name	description	amenities
5	POINT (524555.52...	https://www.airb...	Fulham Garden Flat	Well decorated s...	["Dining
426	POINT (543953.65...	https://www.airb...	Cosy London Flat...	Enjoy my warm an...	["Bed lin
516	POINT (506207.28...	https://www.airb...	Furnished Apartment	Make some memori...	["Kitcher
725	POINT (536247.17...	https://www.airb...	Comfortable room...	My place is clos...	["Hanger
1492	POINT (528407.92...	https://www.airb...	Modern, bright, ...	Luxury canal sid...	["Dining


```
pd.options.display.max_colwidth=75
srcdf[srcdf.Topic==1].description_norm.head(10)
```

```
426      enjoy warm cosy modern apartment . river view gorgeous view city . walk...
516      make memory unique familyfriendly place . near nice bridge minute away ...
725      place close mile tube station  brick lane shoreditch  queen mary univer...
4667     beautiful modern onebedroom flat heart london . covent garden literally...
9867     . bath flat tuck safe quiet neighbourhood  walk tube line . hammersmith...
11117    enjoy trendy experience centrallylocated home . historic bedroom apartm...
12188    make memory unique familyfriendly place . near nice bridge minute away ...
12217    newly renovate  totally equipped furnished modern apartment heart londo...
13405    great newly renovate bedroom apartment second room turn office . perfec...
13429    absolutely enormous room available canary wharf  balcony . room comfort...
Name: description_norm, dtype: object
```

```
vectorizer = CountVectorizer(ngram_range=(1,1), stop_words='english', analyzer='word
topic_corpus = vectorizer.fit_transform(srcdf[srcdf.Topic==1].description.values) #
```

```
topicdf = pd.DataFrame(data=topic_corpus.toarray(),
                        columns=vectorizer.get_feature_names_out())
```

```
f,ax = plt.subplots(1,1,figsize=(5,5))
plt.gcf().set_dpi(300)
Cloud = WordCloud(
    background_color="white",
    max_words=75).generate_from_frequencies(topiccdf.sum())
ax.imshow(Cloud)
ax.axis("off");
# plt.savefig('Wordcloud 3.png')
```



11 Word2Vec

 Tip

This algorithm works almost like magic. You should play with the configuration parameters and see how it changes your results. **However**, at this time I've not preinstalled `gensim` as it bloated the image substantially.

This algorithm works almost like magic. You should play with the configuration parameters and see how it changes your results. **However**, at this time I've not preinstalled `gensim` as it bloated the image substantially.

11.1 Configure

```
from gensim.models.word2vec import Word2Vec
```

```
dims = 100
print(f"You've chosen {dims} dimensions.")
```

```
window = 4
print(f"You've chosen a window of size {window}.")
```

```
min_v_freq = 0.01 # Don't keep words appearing less than 1% frequency
min_v_count = math.ceil(min_v_freq * srcdf.shape[0])
print(f"With a minimum frequency of {min_v_freq} and {srcdf.shape[0]:,}
```

11.2 Train

```
%%time

corpus      = srcdf.description_norm.fillna(' ').values
#corpus_sent = [nltk.sent_tokenize(text) for text in corpus] # <-- with more formal
corpus_sent = [d.replace('.', ' ').split(' ') for d in corpus] # <-- deals better with
model       = Word2Vec(sentences=corpus_sent, vector_size=dims, window=window, epoch
                      min_count=min_v_count, seed=42, workers=1)

#model.save(f"word2vec-d{dims}-w{window}.model") # <-- You can then Word2Vec.load(..
```

11.3 Explore Similarities

This next bit of code only runs if you have calculated the frequencies above in the [Frequencies and Ngrams](#) section.

```
pd.set_option('display.max_colwidth',150)

df = fcounts[1] # <-- copy out only the unigrams as we haven't trained anything else

n      = 14 # number of words
topn   = 7  # number of most similar words

selected_words = df[df['Ngram Size 1'] > 5].reset_index().level_0.sample(n, random_s

words = []
v1     = []
v2     = []
v3     = []
sims   = []

for w in selected_words:
    try:
        vector = model.wv[w] # get numpy vector of a word
        #print(f"Word vector for '{w}' starts: {vector[:5]}...")

        sim = model.wv.most_similar(w, topn=topn)
        #print(f"Similar words to '{w}' include: {sim}.")

        words.append(w)
        v1.append(vector[0])
        v2.append(vector[1])
        v3.append(vector[2])
        sims.append(", ".join([x[0] for x in sim]))
    except KeyError:
        print(f"Didn't find {w} in model. Can happen with low-frequency terms.")

vecs = pd.DataFrame({
```

```

    'Term':words,
    'V1':v1,
    'V2':v2,
    'V3':v3,
    f'Top {topn} Similar':sims
  })

vecs

```

```

#print(model.wv.index_to_key) # <-- the full vocabulary that has been trained

```

11.4 Apply

We're going to make *use* of this further next week...

11.4.1 Questions

- What happens when *dims* is very small (e.g. 25) or very large (e.g. 300)?
- What happens when *window* is very small (e.g. 2) or very large (e.g. 8)?

12 Processing the Full File

Caution

This code can take *some time* (> 5 minutes on a M2 Mac) to run, so **don't run this** until you've understood what we did before!

You will get a warning about "." looks like a filename, not markup — this looks a little scary, but is basically suggesting that we have a description that consists only of a '.' or that looks like some kind of URL (which the parser thinks means you're trying to pass it something to download).

```

%%time
# This can take up to 8 minutes on a M2 Mac
gdf['description_norm'] = ''
gdf['description_norm'] = gdf.description.apply(normalise_document, remove_digits=Tr

gdf.to_parquet(os.path.join('data','geo',f'{fn.replace(".", "-with-nlp.")}'))

```

Tip

Saving an intermediate file at this point is useful because you've done quite a bit of *expensive* computation. You *could* restart-and-run-all and then go out for the day, but probably easier to just save this output and then, if you need to restart your analysis at some point in the future, just remember to deserialise amenities back into a list format.

12.1 Applications

The above is *still* only the results for the one of the subsets of apartments *alone*. At this point, you would probably want to think about how your results might change if you changed any of the following:

1. Using one of the other data sets that we created, or even the entire data set!
2. Applying the CountVectorizer or TfidfVectorizer *before* selecting out any of our ‘sub’ data sets.
3. Using the visualisation of information to improve our regex selection process.
4. Reducing, increasing, or constraining (i.e. `ngrams=(2,2)`) the size of the ngrams while bearing in mind the impact on processing time and interpretability.
5. Filtering by type of listing or host instead of keywords found in the description (for instance, what if you applied TF/IDF to the entire data set and then selected out ‘Whole Properties’ before splitting into those advertised by hosts with only one listing vs. those with multiple listings?).
6. Linking this back to the geography.

Over the next few weeks we’ll also consider alternative means of visualising the data!

12.2 Resources

There is a lot more information out there, including a [whole book](#) and your standard [O’Reilly text](#).

And some more useful links:

- [Pandas String Contains Method](#)
- [Using Regular Expressions with Pandas](#)
- [Summarising Chapters from Frankenstein using TF/IDF](#)