Practical 8bis: Working with Text (Part 2)

The basics of Text Mining and NLP

Table of contents

1 Preamble **2**

2 Setup **2**

3 Illustrative Text Cleaning **6**

4 Applying Normalisation **40**

5 Revenons à Nos Moutons **45**

6 Word Clouds **56**

7 Latent Dirchlet Allocation **58**

8 Word2Vec **61**

9 Processing the Full File **64**

Part 2 of Practical 8 is *optional* and should only be attempted if Part 1 made sense to you.

1. The first few tasks are about finding important vocabulary (think 'keywords' and 'significant terms') in documents so that you can start to think about what is *distinctive* about documents and groups of documents. **This is quite useful and relatively easier to understand than what comes next!**
2. The second part is about fully-fledged NLP using Latent Direclecht Allocation (topic modelling) and Word2Vec (words embeddings for use in clustering or similarity work).

The later parts are largely complete and ready to run; however, that *doesn't* mean you should just skip over them and think you've grasped what's happening and it will be easy to apply in your own analyses. I would *not* pay as much attention to LDA topic mining since I don't think it's results are that good, but I've included it here as it's still commonly-used in the Digital Humanities and by Marketing folks. Word2Vec is much more powerful and forms the basis of the kinds of advances seen in ChatGPT and other LLMs.

> Connections
>
> Working with text is unquestionably *hard*. In fact, *conceptually* this is probaly the most challenging practical of the term! But data scientists are *always* dealing with text because so much of the data that we collect (even more so thanks to the web) is not only text-based (URLs are text!) but, increasingly, unstructured (social media posts, tags, etc.). So while getting to grips with text is a challenge, it also uniquely positions you with respect to the skills and knowledge that other graduates are offering to employers.

1 Preamble

This practical has been written using `nltk`, but would be *relatively* easy to rework using `spacy`. Most programmers tend to use one *or* the other, and the switch wouldn't be hard other than having to first load the requisite language models:

```
import spacy

# `...web_md` and `...web_lg` are also options
corp = "en_core_web_sm"

try:
    nlp = spacy.load(corp)
except OSError:
    spacy.cli.download(corp)
    nlp = spacy.load(corp)
```

You can read about the models, and note that they are also available in other languages besides English.

2 Setup

> Difficulty Level: Low
>
> But this is only because this has been worked out for you. Starting from sctach in NLP is *hard* so people try to avoid it as much as possible.

2.1 Required Modules

> Note
>
> Notice that the number of modules and functions that we import is steadily increasing week-on-week, and that for text processing we tend to draw on quite a wide range of utilies! That said, the three most commonly used are: `sklearn`,

`nltk`, and `spacy`.

Standard libraries we've seen before.

```
import os
import numpy as np
import pandas as pd
import geopandas as gpd
import re
import math
import matplotlib.pyplot as plt
```

Vectorisers we will use from the 'big beast' of Python machine learning: Sci-Kit Learn.

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.decomposition import LatentDirichletAllocation

# We don't use this but I point out where you *could*
from sklearn.preprocessing import OneHotEncoder
```

NLP-specific libraries that we will use for tokenisation, lemmatisation, and frequency analysis.

```
import nltk
import spacy
from nltk.corpus import wordnet as wn
from nltk.stem.wordnet import WordNetLemmatizer

try:
    from nltk.corpus import stopwords
except:
    ! python -c 'import nltk; nltk.download("stopwords"); nltk.download("punkt_tab")
    from nltk.corpus import stopwords

stopword_list = set(stopwords.words('english'))

from nltk.tokenize import word_tokenize, sent_tokenize
from nltk.tokenize.toktok import ToktokTokenizer

from nltk.stem.porter import PorterStemmer
from nltk.stem.snowball import SnowballStemmer

from nltk import ngrams, FreqDist

lemmatizer = WordNetLemmatizer()
tokenizer = ToktokTokenizer()
```

Remaining libraries that we'll use for processing and display text data. Most of this relates to dealing with the various ways that text data cleaning is *hard* because of

the myriad formats it comes in.

```
import string
import unicodedata
from bs4 import BeautifulSoup
from wordcloud import WordCloud, STOPWORDS
```

This next is just a small utility function that allows us to output Markdown (like this cell) instead of plain text:

```
from IPython.display import display_markdown

def as_markdown(head='', body='Some body text'):
    if head != '':
        display_markdown(f"##### {head}\n\n>{body}\n", raw=True)
    else:
        display_markdown(f">{body}\n", raw=True)

as_markdown('Result!', "Here's my output...")
```

Result!

> Here's my output...

2.2 Loading Data

> **Connections**
>
> Because I generally want each practical to stand on its own (unless I'm trying to make a *point*), I've not moved this to a separate Python file (e.g. `utils.py`, but in line with what we covered back in the lectures on Functions and Packages, this sort of thing is a good candidate for being split out to a separate file to simplify re-use.

Remember this function from last week? We use it to save downloading files that we already have stored locally. But notice I've made some small changes... what do these do to help the user?

```
import os
from requests import get
from urllib.parse import urlparse
from functools import wraps

def check_cache(f):
    @wraps(f)
    def wrapper(src, dst, min_size=100):
        url = urlparse(src) # We assume that this is some kind of valid URL
        fn  = os.path.split(url.path)[-1] # Extract the filename
```

```
        dsn = os.path.join(dst,fn) # Destination filename
        if os.path.isfile(dsn) and os.path.getsize(dsn) > min_size:
            print(f"+ {dsn} found locally!")
            return(dsn)
        else:
            print(f"+ {dsn} not found, downloading!")
            return(f(src, dsn))
    return wrapper

@check_cache
def cache_data(src:str, dst:str) -> str:
    """Downloads a remote file.

    The function sits between the 'read' step of a pandas or geopandas
    data frame and downloading the file from a remote location. The idea
    is that it will save it locally so that you don't need to remember to
    do so yourself. Subsequent re-reads of the file will return instantly
    rather than downloading the entire file for a second or n-th itme.

    Parameters
    ----------
    src : str
        The remote *source* for the file, any valid URL should work.
    dst : str
        The *destination* location to save the downloaded file.

    Returns
    -------
    str
        A string representing the local location of the file.
    """

    # Convert the path back into a list (without)
    # the filename -- we need to check that directories
    # exist first.
    path = os.path.split(dst)[0]
    print(f"Path: {path}")

    # Create any missing directories in dest(ination) path
    # -- os.path.join is the reverse of split (as you saw above)
    # but it doesn't work with lists... so I had to google how
    # to use the 'splat' operator! os.makedirs creates missing
    # directories in a path automatically.
    if path != '':
        os.makedirs(path, exist_ok=True)

    # Download and write the file
    with open(dst, "wb") as file:
        response = get(src)
        file.write(response.content)
```

```
    print(' + Done downloading...')

    return dst
```

> **Tip**
>
> For very large *non*-geographic data sets, remember that you can `use_cols` (or `columns` depending on the file type) to specify a subset of columns to load.

Load the main data set:

```
# Load the data sets created in the previous practical
lux    = gpd.read_parquet(os.path.join('data','clean','luxury.geopackage'))
aff    = gpd.read_parquet(os.path.join('data','clean','affordable.geopackage'))
bluesp = gpd.read_parquet(os.path.join('data','clean','bluespace.geopackage'))
```

3 Illustrative Text Cleaning

Now we're going to step through the *parts* of the process that we apply to clean and transform text. We'll do this individually before using a function to apply them *all at once*.

3.1 Downloading a Web Page

> **Difficulty Level: Low.**

There is plenty of good economic geography research being done using web pages. Try using Google Scholar to look for work using the British Library's copy of the *Internet Archive*.

```
from urllib.request import urlopen, Request

# We need this so that the Bartlett web site 'knows'
# what kind of browser it is deasling with. Otherwise
# you get a Permission Error (403 Forbidden) because
# the site doesn't know what to do.
hdrs = {
    'User-Agent': 'Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/537.11 (KHTML, like G
    'Accept': 'text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8',
}
url = 'https://www.ucl.ac.uk/bartlett/casa/about-0'
```

Question

```
# Notice that here we have to assemble a request and
# then 'open' it so that the request is properly issued
# to the web server. Normally, we'd just use `urlopen`,
# but that doesn't give you the ability to set the headers.
request  = Request(url, None, hdrs) #The assembled request
response = urlopen(request)
html     = response.???.decode('utf-8') # The data you need

print(html[:1000])
```

Answer

```
# Notice that here we have to assemble a request and
# then 'open' it so that the request is properly issued
# to the web server. Normally, we'd just use `urlopen`,
# but that doesn't give you the ability to set the headers.
request  = Request(url, None, hdrs) #The assembled request
response = urlopen(request)
html     = response.read().decode('utf-8') # The data you need

print(html[:1000])
```

```
<!DOCTYPE html>
<html lang="en" dir="ltr" prefix="og: https://ogp.me/ns#">
  <head>
    <script>(function() {
  'use strict';

  // Get whitelisted domains from PHP parameter.
  const whitelistedDomains = ["https:\/\/app.geckoform.com"];

  // Convert whitelisted domains to Set for O(1) lookup performance.
  const whitelistedSet = new Set(whitelistedDomains);

  /**
   * Check if a URL should be whitelisted.
   * @param {string} url - The URL to check.
   * @returns {boolean} - True if the URL should be whitelisted.
   */
  const isWhitelisted = (url) => {
    if (!url) return false;
    for (const domain of whitelistedSet) {
      if (url.startsWith(domain)) return true;
    }
    return false;
  };

  /**
   * This function finds an iframe, moves its src to data-src, and sets
```

```
 * the src to about:blank to stop it from loading.
* This prevents third-party cookies from being set before consent is given.
 */
const preventLoad = (iframe) => {
  // Early return if iframe is alre
```

3.2 Removing HTML

> ⚠ Difficulty level: Moderate
>
> Because what we're doing will seem really strange and uses some previously unseen libraries that you'll have to google.

Hint: you need to need to **get the text** out of the each returned `<p>` and `<div>` element! I'd suggest also commenting this up since there is a *lot* going on on some of these lines of code!

Question

```
cleaned = []

soup = BeautifulSoup(html)
body = soup.find('body')

for c in body.findChildren(recursive=False):
    if c.name in ['div','p'] and c.???.strip() != '':
        # \xa0 is a non-breaking space in Unicode (  in HTML)
        txt = [re.sub(r'(?:\u202f|\xa0|\u200b)',' ',x.strip()) for x in c.get_text(s
        cleaned += txt

cleaned
```

Answer

```
cleaned = []

soup = BeautifulSoup(html)
body = soup.find('body')

for c in body.findChildren(recursive=False):
    if c.name in ['div','p'] and c.get_text().strip() != '':
        # \xa0 is a non-breaking space in Unicode (  in HTML)
        txt = [re.sub(r'(?:\u202f|\xa0|\u200b)',' ',x.strip()) for x in c.get_text(s
        cleaned += txt

cleaned
```

```
['Navigate back to homepage',
 'Open search bar.',
 'Open main navigation menu',
 'Search Term',
 '(optional)',
 'Submit search',
 'Main navigation',
 'Study',
 'Study at UCL',
 'Being a student at UCL is about so much more than just acquiring knowledge. Studying here g
 'Undergraduate courses',
 'Graduate courses',
 'Short courses',
 'Study abroad',
 'Centre for Languages & International Education',
 'Close sub menu modal',
 'Research',
 'Research at UCL',
 'Find out more about what makes UCL research world-leading, how to access UCL expertise, ar
 'Engage with us',
 'Explore our Research',
 'Initiatives and networks',
 'Research news',
 'Close sub menu modal',
 'Engage',
 'Engage with UCL',
 'Discover the many ways you can connect with UCL, and how we work with industry, government
 'Alumni',
 'Business partnerships and collaboration',
 'Global engagement',
 'News and Media relations',
 'Public Policy',
 'Schools and priority groups',
 'Give to UCL',
 'Close sub menu modal',
 'About',
 'About UCL',
 "Founded in 1826 in the heart of London, UCL is London's leading multidisciplinary univers
 'Who we are',
 'Faculties',
 'Governance',
 'President and Provost',
 'Strategy',
 'Close sub menu modal',
 'Active parent page:',
 'The Bartlett Faculty of the Built Environment',
 'Study',
 'Research',
 'Active parent page:',
 'Our schools and institutes',
 'People',
 'Ideas',
```

```
 'Engage',
 'News and Events',
 'About',
 'Open or close faculty menu',
 'Breadcrumb trail',
 'The Bartlett Faculty of the Built Environment',
 'Our schools and institutes',
 'The Bartlett Centre for Advanced Spatial Analysis',
 'Faculty menu',
 'Current page:',
 'About',
 'CUPUM 2025',
 'Partnerships',
 'Software',
 'Remembering Sinesio Alves Junior',
 'About',
 'The Centre for Advanced Spatial Analysis (CASA) is an interdisciplinary research institut
 'The Centre for Advanced Spatial Analysis (CASA) was established in 1995 to lead the develo
including physicists, planners, geographers, economists, data scientists, architects, math
united by our mission to tackle the biggest challenges facing cities and societies around th
 'Contact us',
 '90 Tottenham Court Road',
 'London',
 'W1T 4TJ',
 'UCL Centre for Advanced Spatial Analysis',
 'Click to email.',
 'casa@ucl.ac.uk',
 'Click to call.',
 '+44 (0)20 3108 4910',
 'UCL footer',
 'Visit',
 'Bloomsbury Theatre and Studio',
 'Library, Museums and Collections',
 'UCL Maps',
 'UCL Shop',
 'Contact UCL',
 'Students',
 'Accommodation',
 'Current Students',
 'Moodle',
 "Students' Union",
 'Staff',
 'Inside UCL',
 'Staff Intranet',
 'Work at UCL',
 'Human Resources',
 'UCL social media menu',
 'Link to Soundcloud',
 'Link to Flickr',
 'Link to TikTok',
 'Link to Youtube',
 'Link to Instagram',
```

```
 'Link to Facebook',
 'Link to Twitter',
 'University College London, Gower Street, London, WC1E 6BT',
 'Tel: +44 (0) 20 7679 2000',
 '© 2025 UCL',
 'Essential',
 'Disclaimer',
 'Freedom of Information',
 'Accessibility',
 'Cookies',
 'Privacy',
 'Slavery statement',
 'Log in']
```

3.3 Lower Case

> Difficulty Level: Low.

Question

```
lower = [c.???() for ??? in cleaned]
lower
```

Answer

```
lower = [s.lower() for s in cleaned]
lower
```

```
['navigate back to homepage',
 'open search bar.',
 'open main navigation menu',
 'search term',
 '(optional)',
 'submit search',
 'main navigation',
 'study',
 'study at ucl',
 'being a student at ucl is about so much more than just acquiring knowledge. studying here g
 'undergraduate courses',
 'graduate courses',
 'short courses',
 'study abroad',
 'centre for languages & international education',
 'close sub menu modal',
 'research',
 'research at ucl',
```

```
'find out more about what makes ucl research world-leading, how to access ucl expertise, ar
'engage with us',
'explore our research',
'initiatives and networks',
'research news',
'close sub menu modal',
'engage',
'engage with ucl',
'discover the many ways you can connect with ucl, and how we work with industry, government
'alumni',
'business partnerships and collaboration',
'global engagement',
'news and media relations',
'public policy',
'schools and priority groups',
'give to ucl',
'close sub menu modal',
'about',
'about ucl',
"founded in 1826 in the heart of london, ucl is london's leading multidisciplinary univers
'who we are',
'faculties',
'governance',
'president and provost',
'strategy',
'close sub menu modal',
'active parent page:',
'the bartlett faculty of the built environment',
'study',
'research',
'active parent page:',
'our schools and institutes',
'people',
'ideas',
'engage',
'news and events',
'about',
'open or close faculty menu',
'breadcrumb trail',
'the bartlett faculty of the built environment',
'our schools and institutes',
'the bartlett centre for advanced spatial analysis',
'faculty menu',
'current page:',
'about',
'cupum 2025',
'partnerships',
'software',
'remembering sinesio alves junior',
'about',
'the centre for advanced spatial analysis (casa) is an interdisciplinary research institut
'the centre for advanced spatial analysis (casa) was established in 1995 to lead the develo
```

including physicists, planners, geographers, economists, data scientists, architects, math
united by our mission to tackle the biggest challenges facing cities and societies around th
 'contact us',
 '90 tottenham court road',
 'london',
 'w1t 4tj',
 'ucl centre for advanced spatial analysis',
 'click to email.',
 'casa@ucl.ac.uk',
 'click to call.',
 '+44 (0)20 3108 4910',
 'ucl footer',
 'visit',
 'bloomsbury theatre and studio',
 'library, museums and collections',
 'ucl maps',
 'ucl shop',
 'contact ucl',
 'students',
 'accommodation',
 'current students',
 'moodle',
 "students' union",
 'staff',
 'inside ucl',
 'staff intranet',
 'work at ucl',
 'human resources',
 'ucl social media menu',
 'link to soundcloud',
 'link to flickr',
 'link to tiktok',
 'link to youtube',
 'link to instagram',
 'link to facebook',
 'link to twitter',
 'university college london, gower street, london, wc1e 6bt',
 'tel: +44 (0) 20 7679 2000',
 '© 2025 ucl',
 'essential',
 'disclaimer',
 'freedom of information',
 'accessibility',
 'cookies',
 'privacy',
 'slavery statement',
 'log in']

3.4 Stripping 'Punctuation'

> Difficulty level: Hard
>
> This is because you need to understand: 1) why we're *compiling* the regular expression and how to use character classes; and 2) how the NLTK tokenizer differs in approach to the regex.

3.4.1 Regular Expression Approach

We want to clear out punctuation using a regex that takes advantage of the `[...]` (character class) syntax. The really tricky part is remembering how to specify the 'punctuation' when some of that punctuation has 'special' meanings in a regular expression context. For instance, `.` means 'any character', while `[` and `]` mean 'character class'. So this is another *escaping* problem and it works the *same* way it did when we were dealing with the Terminal...

Hints: some other factors...

1. You will want to match more than one piece of punctuation at a time, so I'd suggest add a `+` to your pattern.
2. You will need to look into *metacharacters* for creating a kind of 'any of the characters *in this class*' bag of possible matches.

Question

```
pattern = re.compile(r'[???]+')
print(pattern)
```

Answer

```
pattern = re.compile(r'[,\.!\-><=\(\)\[\]\/&\'\"';\+\-\—]+')
print(pattern)
```

```
re.compile('[,\\.!\\-><=\\(\\)\\[\\]\\/&\\\'\\"';\\+\\-\\—]+')
```

3.4.2 Tokenizer

The other way to do this, which is probably *easier* but produces more complex output, is to draw on the tokenizers already provided by NLTK. For our purposes `word_tokenize` is probably fine, but depending on your needs there are other options and you can also write your own.

```
nltk.download('punkt')
nltk.download('wordnet')
from nltk.tokenize import word_tokenize
```

```
print(word_tokenize)
```

```
<function word_tokenize at 0x314cb7600>
```

```
[nltk_data] Downloading package punkt to /Users/jreades/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /Users/jreades/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

3.4.3 Compare

Look at how these outputs differ in subtle ways:

```
subbed = []
tokens = []
for l in lower:
    subbed.append(re.sub(pattern, ' ', l))
    tokens.append(word_tokenize(l))

for s in subbed:
    as_markdown("Substituted", s)

for t in tokens:
    as_markdown("Tokenised", t)
```

Substituted

navigate back to homepage

Substituted

open search bar

Substituted

open main navigation menu

Substituted

search term

Substituted

optional

Substituted

submit search

Substituted

main navigation

Substituted

study

Substituted

study at ucl

Substituted

being a student at ucl is about so much more than just acquiring knowledge studying here gives you the opportunity to realise your potential as an individual and the skills and tools to thrive

Substituted

undergraduate courses

Substituted

graduate courses

Substituted

short courses

Substituted

study abroad

Substituted

centre for languages international education

Substituted

close sub menu modal

Substituted

research

Substituted

research at ucl

Substituted

find out more about what makes ucl research world leading how to access ucl expertise and teams in the office of the vice provost research innovation and global engagement

Substituted

engage with us

Substituted

explore our research

Substituted

initiatives and networks

Substituted

research news

Substituted

close sub menu modal

Substituted

engage

Substituted

engage with ucl

Substituted

discover the many ways you can connect with ucl and how we work with industry government and not for profit organisations to tackle tough challenges

Substituted

alumni

Substituted

business partnerships and collaboration

Substituted

global engagement

Substituted

news and media relations

Substituted

public policy

Substituted

schools and priority groups

Substituted

give to ucl

Substituted

close sub menu modal

Substituted

about

Substituted

about ucl

Substituted

founded in 1826 in the heart of london ucl is london s leading multidisciplinary university with more than 16 000 staff and 50 000 students from 150 different countries

Substituted

who we are

Substituted

faculties

Substituted

governance

Substituted

president and provost

Substituted

strategy

Substituted

close sub menu modal

Substituted

active parent page:

Substituted

the bartlett faculty of the built environment

Substituted

study

Substituted

research

Substituted

active parent page:

Substituted

our schools and institutes

Substituted

people

Substituted

ideas

Substituted

engage

Substituted

news and events

Substituted

about

Substituted

open or close faculty menu

Substituted

breadcrumb trail

Substituted

the bartlett faculty of the built environment

Substituted

our schools and institutes

Substituted

the bartlett centre for advanced spatial analysis

Substituted

faculty menu

Substituted

current page:

Substituted

about

Substituted

cupum 2025

Substituted

partnerships

Substituted

software

Substituted

remembering sinesio alves junior

Substituted

about

Substituted

the centre for advanced spatial analysis casa is an interdisciplinary research institute focusing on the science of cities within the bartlett faculty of the built environment at ucl

Substituted

the centre for advanced spatial analysis casa was established in 1995 to lead the development of a science of cities drawing upon methods and ideas in modelling and data science sensing the urban environment visualisation and computation advancing urban science through research and collaboration today casa s research is still pushing boundaries to create better cities for everyone both leading the intellectual agenda and working closely with government and industry partners to make real world impact our teaching reflects this making the most of our cutting edge research tools new forms of data and long standing non academic partnerships to train the next generation of urban scientists with the skills and ideas they ll need to have an impact in industry government academia and the third sector the casa community is closely connected but strongly interdisciplinary we bring together people from around the world with a unique variety of backgrounds including physicists planners geographers economists data scientists architects mathematicians and computer scientists united by our mission to tackle the biggest challenges facing cities and societies around the world we work across multiple scales: from the hyper local environment of the low powered sensor all the way up to satellite remote sensing of whole countries and regions studying at casa studying at casa brings lifelong value with our students poised to take on leadership and integration roles at the forefront of urban and spatial data science by studying with us you will become part of our active and engaged alumni community with access to job listings networking and social activities as well as continued contact with our outstanding teachers and researchers find out more about connected environments at casa working at casa we employ around 40 permanent academic professional services and research staff and our researchers students and alumni work at the cutting edge of this interdisciplinary science our research tools and models frequently inform systems ranging from urban planning to museum retail and heritage applications and onwards to social network understanding vacancies in the department will be published on the ucl jobs page connect with us on linkedin

Substituted

contact us

Substituted

90 tottenham court road

Substituted

london

Substituted

w1t 4tj

Substituted

ucl centre for advanced spatial analysis

Substituted

click to email

Substituted

casa@ucl ac uk

Substituted

click to call

Substituted

44 0 20 3108 4910

Substituted

ucl footer

Substituted

visit

Substituted

bloomsbury theatre and studio

Substituted

library museums and collections

Substituted

ucl maps

Substituted

ucl shop

Substituted

contact ucl

Substituted

students

Substituted

accommodation

Substituted

current students

Substituted

moodle

Substituted

students union

Substituted

staff

Substituted

inside ucl

Substituted

staff intranet

Substituted

work at ucl

Substituted

human resources

Substituted

ucl social media menu

Substituted

link to soundcloud

Substituted

link to flickr

Substituted

link to tiktok

Substituted

link to youtube

Substituted

link to instagram

Substituted

link to facebook

Substituted

link to twitter

Substituted

university college london gower street london wc1e 6bt

Substituted

tel: 44 0 20 7679 2000

Substituted

© 2025 ucl

Substituted

essential

Substituted

disclaimer

Substituted

freedom of information

Substituted

accessibility

Substituted

cookies

Substituted

privacy

Substituted

slavery statement

Substituted

log in

Tokenised

[‘navigate’, ‘back’, ‘to’, ‘homepage’]

Tokenised

[‘open’, ‘search’, ‘bar’, ‘:’]

Tokenised

[‘open’, ‘main’, ‘navigation’, ‘menu’]

Tokenised

['search', 'term']

Tokenised

['(', 'optional', ')']

Tokenised

['submit', 'search']

Tokenised

['main', 'navigation']

Tokenised

['study']

Tokenised

['study', 'at', 'ucl']

Tokenised

['being', 'a', 'student', 'at', 'ucl', 'is', 'about', 'so', 'much', 'more', 'than', 'just', 'acquiring', 'knowledge', '.', 'studying', 'here', 'gives', 'you', 'the', 'opportunity', 'to', 'realise', 'your', 'potential', 'as', 'an', 'individual', ',', 'and', 'the', 'skills', 'and', 'tools', 'to', 'thrive', '.']

Tokenised

['undergraduate', 'courses']

Tokenised

['graduate', 'courses']

Tokenised

['short', 'courses']

Tokenised

['study', 'abroad']

Tokenised

['centre', 'for', 'languages', '&', 'international', 'education']

Tokenised

['close', 'sub', 'menu', 'modal']

Tokenised

['research']

Tokenised

['research', 'at', 'ucl']

Tokenised

['find', 'out', 'more', 'about', 'what', 'makes', 'ucl', 'research', 'world-leading', ',', 'how', 'to', 'access', 'ucl', 'expertise', ',', 'and', 'teams', 'in', 'the', 'office', 'of', 'the', 'vice-provost', '(', 'research', ',', 'innovation', 'and', 'global', 'engagement', ')', '.']

Tokenised

['engage', 'with', 'us']

Tokenised

['explore', 'our', 'research']

Tokenised

['initiatives', 'and', 'networks']

Tokenised

['research', 'news']

Tokenised

['close', 'sub', 'menu', 'modal']

Tokenised

['engage']

Tokenised

['engage', 'with', 'ucl']

Tokenised

['discover', 'the', 'many', 'ways', 'you', 'can', 'connect', 'with', 'ucl', ',', 'and', 'how', 'we', 'work', 'with', 'industry', ',', 'government', 'and', 'not-for-profit', 'organisations', 'to', 'tackle', 'tough', 'challenges', '.']

Tokenised

['alumni']

Tokenised

['business', 'partnerships', 'and', 'collaboration']

Tokenised

['global', 'engagement']

Tokenised

['news', 'and', 'media', 'relations']

Tokenised

['public', 'policy']

Tokenised

['schools', 'and', 'priority', 'groups']

Tokenised

['give', 'to', 'ucl']

Tokenised

['close', 'sub', 'menu', 'modal']

Tokenised

['about']

Tokenised

['about', 'ucl']

Tokenised

['founded', 'in', '1826', 'in', 'the', 'heart', 'of', 'london', ',', 'ucl', 'is', 'london', "'s", 'leading', 'multidisciplinary', 'university', ',', 'with', 'more', 'than', '16,000', 'staff', 'and', '50,000', 'students', 'from', '150', 'different', 'countries', '.']

Tokenised

['who', 'we', 'are']

Tokenised

['faculties']

Tokenised

['governance']

Tokenised

['president', 'and', 'provost']

Tokenised

['strategy']

Tokenised

['close', 'sub', 'menu', 'modal']

Tokenised

['active', 'parent', 'page', ':']

Tokenised

['the', 'bartlett', 'faculty', 'of', 'the', 'built', 'environment']

Tokenised

['study']

Tokenised

['research']

Tokenised

['active', 'parent', 'page', ':']

Tokenised

['our', 'schools', 'and', 'institutes']

Tokenised

['people']

Tokenised

['ideas']

Tokenised

['engage']

Tokenised

['news', 'and', 'events']

Tokenised

['about']

Tokenised

['open', 'or', 'close', 'faculty', 'menu']

Tokenised

['breadcrumb', 'trail']

Tokenised

['the', 'bartlett', 'faculty', 'of', 'the', 'built', 'environment']

Tokenised

['our', 'schools', 'and', 'institutes']

Tokenised

['the', 'bartlett', 'centre', 'for', 'advanced', 'spatial', 'analysis']

Tokenised

['faculty', 'menu']

Tokenised

['current', 'page', ':']

Tokenised

['about']

Tokenised

['cupum', '2025']

Tokenised

['partnerships']

Tokenised

['software']

Tokenised

['remembering', 'sinesio', 'alves', 'junior']

Tokenised

['about']

Tokenised

['the', 'centre', 'for', 'advanced', 'spatial', 'analysis', '(', 'casa', ')', 'is', 'an', 'interdisciplinary', 'research', 'institute', 'focusing', 'on', 'the', 'science', 'of', 'cities', 'within', 'the', 'bartlett', 'faculty', 'of', 'the', 'built', 'environment', 'at', 'ucl', '.']

Tokenised

['the', 'centre', 'for', 'advanced', 'spatial', 'analysis', '(', 'casa', ')', 'was', 'established', 'in', '1995', 'to', 'lead', 'the', 'development', 'of', 'a', 'science', 'of', 'cities', 'drawing', 'upon', 'methods', 'and', 'ideas', 'in', 'modelling', 'and', 'data', 'science', ',', 'sensing', 'the', 'urban', 'environment', ',', 'visualisation', 'and', 'computation', '.', 'advancing', 'urban', 'science', 'through', 'research', 'and', 'collaboration', 'today', ',', 'casa', '’', 's', 'research', 'is', 'still', 'pushing', 'boundaries', 'to', 'create', 'better', 'cities', 'for', 'everyone', ',', 'both', 'leading', 'the', 'intellectual', 'agenda', 'and', 'working', 'closely', 'with', 'government', 'and', 'industry', 'partners', 'to', 'make', 'real-world', 'impact', '.', 'our', 'teaching', 'reflects', 'this', ',', 'making', 'the', 'most', 'of', 'our', 'cutting-edge', 'research', ',', 'tools', ',', 'new', 'forms', 'of', 'data', ',', 'and', 'long-standing', 'non-academic', 'partnerships', 'to', 'train', 'the', 'next', 'generation', 'of', 'urban', 'scientists', 'with', 'the', 'skills', 'and', 'ideas', 'they', '’', 'll', 'need', 'to', 'have', 'an', 'impact', 'in', 'industry', ',', 'government', ',', 'academia', ',', 'and', 'the', 'third', 'sector', '.', 'the', 'casa', 'community', 'is', 'closely', 'connected', ',', 'but', 'strongly', 'interdisciplinary', '.', 'we', 'bring', 'together', 'people', 'from', 'around', 'the', 'world', 'with', 'a', 'unique', 'variety', 'of', 'backgrounds', '–', 'including', 'physicists', ',', 'planners', ',', 'geographers', ',', 'economists', ',', 'data', 'scientists', ',', 'architects', ',', 'mathematicians', 'and', 'computer', 'scientists', '–', 'united', 'by', 'our', 'mission', 'to', 'tackle', 'the', 'biggest', 'challenges', 'facing', 'cities', 'and', 'societies', 'around', 'the', 'world', '.', 'we', 'work', 'across', 'multiple', 'scales', ':', 'from', 'the', 'hyper-local', 'environment', 'of', 'the', 'low-powered', 'sensor', 'all', 'the', 'way', 'up', 'to', 'satellite', 'remote', 'sensing', 'of', 'whole', 'countries', 'and', 'regions', '.', 'studying', 'at', 'casa', 'studying', 'at', 'casa', 'brings', 'lifelong', 'value', ',', 'with', 'our', 'students', 'poised', 'to', 'take', 'on', 'leadership', 'and', 'integration', 'roles', 'at', 'the', 'forefront', 'of', 'urban', 'and', 'spatial', 'data', 'science', '.', 'by', 'studying', 'with', 'us', 'you', 'will', 'become', 'part', 'of', 'our', 'active', 'and', 'engaged', 'alumni', 'community', ',', 'with', 'access', 'to', 'job', 'listings', ',', 'networking', 'and', 'social', 'activities', ',', 'as', 'well', 'as', 'continued', 'contact', 'with', 'our', 'outstanding', 'teachers', 'and', 'researchers', '.', 'find', 'out', 'more', 'about', 'connected', 'environments', 'at', 'casa', 'working', 'at', 'casa', 'we', 'employ', 'around', '40', 'permanent', 'academic', ',', 'professional', 'services', 'and', 'research', 'staff', ',', 'and', 'our', 'researchers', ',', 'students', 'and', 'alumni', 'work', 'at', 'the', 'cutting', 'edge', 'of', 'this', 'interdisciplinary', 'science', '.', 'our', 'research', 'tools', 'and', 'models', 'frequently', 'inform', 'systems', 'ranging', 'from', 'urban', 'planning', 'to', 'museum', ',', 'retail', ',', 'and', 'heritage', 'applications', 'and', 'onwards', 'to', 'social', 'network', 'understanding', '.', 'vacancies', 'in', 'the', 'department', 'will', 'be', 'published', 'on', 'the', 'ucl', 'jobs', 'page', '.', 'connect', 'with', 'us', 'on', 'linkedin']

Tokenised

['contact', 'us']

Tokenised

['90', 'tottenham', 'court', 'road']

Tokenised

['london']

Tokenised

['w1t', '4tj']

Tokenised

['ucl', 'centre', 'for', 'advanced', 'spatial', 'analysis']

Tokenised

['click', 'to', 'email', ':']

Tokenised

['casa', '@', 'ucl.ac.uk']

Tokenised

['click', 'to', 'call', ':']

Tokenised

['+44', '(', '0', ')', '20', '3108', '4910']

Tokenised

['ucl', 'footer']

Tokenised

['visit']

Tokenised

['bloomsbury', 'theatre', 'and', 'studio']

Tokenised

['library', ',', 'museums', 'and', 'collections']

Tokenised

['ucl', 'maps']

Tokenised

['ucl', 'shop']

Tokenised

['contact', 'ucl']

Tokenised

['students']

Tokenised

['accommodation']

Tokenised

['current', 'students']

Tokenised

['moodle']

Tokenised

['students', " "', 'union']

Tokenised

['staff']

Tokenised

['inside', 'ucl']

Tokenised

['staff', 'intranet']

Tokenised

['work', 'at', 'ucl']

Tokenised

['human', 'resources']

Tokenised

['ucl', 'social', 'media', 'menu']

Tokenised

['link', 'to', 'soundcloud']

Tokenised

['link', 'to', 'flickr']

Tokenised

['link', 'to', 'tiktok']

Tokenised

['link', 'to', 'youtube']

Tokenised

['link', 'to', 'instagram']

Tokenised

['link', 'to', 'facebook']

Tokenised

['link', 'to', 'twitter']

Tokenised

['university', 'college', 'london', ',', 'gower', 'street', ',', 'london', ',', 'wc1e', '6bt']

Tokenised

['tel', ':', '+44', '(', '0', ')', '20', '7679', '2000']

Tokenised

['©', '2025', 'ucl']

Tokenised

['essential']

Tokenised

['disclaimer']

Tokenised

['freedom', 'of', 'information']

Tokenised

['accessibility']

Tokenised

['cookies']

Tokenised

['privacy']

Tokenised

['slavery', 'statement']

Tokenised

['log', 'in']

3.5 Stopword Removal

> ⚠ Difficulty Level: Moderate
>
> You need to remember how list comprehensions work to use the `stopword_list`.

```
stopword_list = set(stopwords.words('english'))
print(stopword_list)
```

```
{'my', 'now', 'wasn', 'shan', 'our', "weren't", 'between', 'needn', 'up', 'not', 'few', 'th
```

Question

```
stopped = []
for p in tokens[2:4]: # <-- why do I just take these items from the list?
    stopped.append([x for x in p if x not in ??? and len(x) > 1])

for s in stopped:
    as_markdown("Line", s)
```

Answer

```
stopped = []
for p in tokens[2:4]: # <-- why do I just take these items from the list?
    stopped.append([x for x in p if x not in stopword_list and len(x) > 1])

for s in stopped:
    as_markdown("Line", s)
```

Line

['open', 'main', 'navigation', 'menu']

Line

['search', 'term']

3.6 Lemmatisation vs Stemming

> 💡 Difficulty level: Low.

```
from nltk.stem.porter import PorterStemmer
from nltk.stem.snowball import SnowballStemmer
```

```
from nltk.stem.wordnet import WordNetLemmatizer
```

```
lemmatizer = WordNetLemmatizer()
print(lemmatizer.lemmatize('monkeys'))
print(lemmatizer.lemmatize('cities'))
print(lemmatizer.lemmatize('complexity'))
print(lemmatizer.lemmatize('Reades'))
```

```
monkey
city
complexity
Reades
```

```
stemmer = PorterStemmer()
print(stemmer.stem('monkeys'))
print(stemmer.stem('cities'))
print(stemmer.stem('complexity'))
print(stemmer.stem('Reades'))
```

```
monkey
citi
complex
read
```

```
stemmer = SnowballStemmer(language='english')
print(stemmer.stem('monkeys'))
print(stemmer.stem('cities'))
print(stemmer.stem('complexity'))
print(stemmer.stem('Reades'))
```

```
monkey
citi
complex
read
```

```
lemmatizer = WordNetLemmatizer()
lemmas  = []
stemmed = []

# This would be better if we passed in a PoS (Part of Speech) tag as well,
# but processing text for parts of speech is *expensive* and for the purposes
# of this tutorial, not necessary.
for s in stopped:
    lemmas.append([lemmatizer.lemmatize(x) for x in s])

for s in stopped:
    stemmed.append([stemmer.stem(x) for x in s])
```

```
for l in lemmas:
    as_markdown('Lemmatised',l)

for s in stemmed:
    as_markdown('Stemmed',s)
```

Lemmatised

['open', 'main', 'navigation', 'menu']

Lemmatised

['search', 'term']

Stemmed

['open', 'main', 'navig', 'menu']

Stemmed

['search', 'term']

```
# What are we doing here?
for ix, p in enumerate(stopped):
    stopped_set = set(stopped[ix])
    lemma_set   = set(lemmas[ix])
    print(sorted(stopped_set.symmetric_difference(lemma_set)))
```

```
[]
[]
```

4 Applying Normalisation

The above approach is fairly hard going since you need to loop through every list element applying these changes one at a time. Instead, we could convert the column to a corpus (or use pandas `apply`) together with a function imported from a library to do the work.

4.1 Downloading the Custom Module

> Difficulty level: Low.

This custom module is not perfect, but it gets the job done... mostly and has some additional features that you could play around with for a final project (e.g. `detect_entities` and `detect_acronyms`).

```
import urllib.request
host  = 'https://orca.casa.ucl.ac.uk'
turl  = f'{host}/~jreades/__textual__.py'
tdirs = os.path.join('textual')
tpath = os.path.join(tdirs,'__init__.py')

if not os.path.exists(tpath):
    os.makedirs(tdirs, exist_ok=True)
    urllib.request.urlretrieve(turl, tpath)
```

4.2 Importing the Custom Module

> Difficulty Level: Low.
>
> But only because you didn't have to write the module! However, the questions could be hard...

In a Jupyter notebook, this code allows us to edit and reload the library dynamically:

```
%load_ext autoreload
%autoreload 2
```

Now let's import it.

```
from textual import *
```

```
All NLTK libraries installed...
```

```
as_markdown('Input', cleaned)
```

Input

['Navigate back to homepage', 'Open search bar.', 'Open main navigation menu', 'Search Term', '(optional)', 'Submit search', 'Main navigation', 'Study', 'Study at UCL', 'Being a student at UCL is about so much more than just acquiring knowledge. Studying here gives you the opportunity to realise your potential as an individual, and the skills and tools to

thrive.', 'Undergraduate courses', 'Graduate courses', 'Short courses', 'Study abroad', 'Centre for Languages & International Education', 'Close sub menu modal', 'Research', 'Research at UCL', 'Find out more about what makes UCL research world-leading, how to access UCL expertise, and teams in the Office of the Vice-Provost (Research, Innovation and Global Engagement).', 'Engage with us', 'Explore our Research', 'Initiatives and networks', 'Research news', 'Close sub menu modal', 'Engage', 'Engage with UCL', 'Discover the many ways you can connect with UCL, and how we work with industry, government and not-for-profit organisations to tackle tough challenges.', 'Alumni', 'Business partnerships and collaboration', 'Global engagement', 'News and Media relations', 'Public Policy', 'Schools and priority groups', 'Give to UCL', 'Close sub menu modal', 'About', 'About UCL', "Founded in 1826 in the heart of London, UCL is London's leading multidisciplinary university, with more than 16,000 staff and 50,000 students from 150 different countries.", 'Who we are', 'Faculties', 'Governance', 'President and Provost', 'Strategy', 'Close sub menu modal', 'Active parent page:', 'The Bartlett Faculty of the Built Environment', 'Study', 'Research', 'Active parent page:', 'Our schools and institutes', 'People', 'Ideas', 'Engage', 'News and Events', 'About', 'Open or close faculty menu', 'Breadcrumb trail', 'The Bartlett Faculty of the Built Environment', 'Our schools and institutes', 'The Bartlett Centre for Advanced Spatial Analysis', 'Faculty menu', 'Current page:', 'About', 'CUPUM 2025', 'Partnerships', 'Software', 'Remembering Sinesio Alves Junior', 'About', 'The Centre for Advanced Spatial Analysis (CASA) is an interdisciplinary research institute focusing on the science of cities within The Bartlett Faculty of the Built Environment at UCL.', 'The Centre for Advanced Spatial Analysis (CASA) was established in 1995 to lead the development of a science of cities drawing upon methods and ideas in modelling and data science, sensing the urban environment, visualisation and computation. Advancing Urban Science through research and collaboration Today, CASA's research is still pushing boundaries to create better cities for everyone, both leading the intellectual agenda and working closely with government and industry partners to make real-world impact. Our teaching reflects this, making the most of our cutting-edge research, tools, new forms of data, and long-standing non-academic partnerships to train the next generation of urban scientists with the skills and ideas they'll need to have an impact in industry, government, academia, and the third sector. The CASA community is closely connected, but strongly interdisciplinary. We bring together people from around the world with a unique variety of backgrounds – including physicists, planners, geographers, economists, data scientists, architects, mathematicians and computer scientists – united by our mission to tackle the biggest challenges facing cities and societies around the world. We work across multiple scales: from the hyper-local environment of the low-powered sensor all the way up to satellite remote sensing of whole countries and regions. Studying at CASA Studying at CASA brings lifelong value, with our students poised to take on leadership and integration roles at the forefront of urban and spatial data science. By studying with us you will become part of our active and engaged alumni community, with access to job listings, networking and social activities, as well as continued contact with our outstanding teachers and researchers. Find out more about Connected Environments at CASA Working at CASA We employ

around 40 permanent academic, professional services and research staff, and our researchers, students and alumni work at the cutting edge of this interdisciplinary science. Our research tools and models frequently inform systems ranging from urban planning to museum, retail, and heritage applications and onwards to social network understanding. Vacancies in the department will be published on the UCL Jobs page . Connect with us on LinkedIn', 'Contact us', '90 Tottenham Court Road', 'London', 'W1T 4TJ', 'UCL Centre for Advanced Spatial Analysis', 'Click to email.', 'casa@ucl.ac.uk', 'Click to call.', '+44 (0)20 3108 4910', 'UCL footer', 'Visit', 'Bloomsbury Theatre and Studio', 'Library, Museums and Collections', 'UCL Maps', 'UCL Shop', 'Contact UCL', 'Students', 'Accommodation', 'Current Students', 'Moodle', "Students' Union", 'Staff', 'Inside UCL', 'Staff Intranet', 'Work at UCL', 'Human Resources', 'UCL social media menu', 'Link to Soundcloud', 'Link to Flickr', 'Link to TikTok', 'Link to Youtube', 'Link to Instagram', 'Link to Facebook', 'Link to Twitter', 'University College London, Gower Street, London, WC1E 6BT', 'Tel: +44 (0) 20 7679 2000', '© 2025 UCL', 'Essential', 'Disclaimer', 'Freedom of Information', 'Accessibility', 'Cookies', 'Privacy', 'Slavery statement', 'Log in']

```
as_markdown('Normalised', [normalise_document(x, remove_digits=True) for x in cleane
```

Normalised

['navigate back homepage', 'open search .', 'open main navigation menu', 'search term', '. optional .', 'submit search', 'main navigation', 'study', 'study', 'student much acquire knowledge . studying give opportunity realise potential individual skill tool thrive .', 'undergraduate course', 'graduate course', 'short course', 'study abroad', 'centre languages international education', 'close menu modal', 'research', 'research', 'find make research worldleading access expertise team office viceprovost . research innovation global engagement .', 'engage', 'explore research', 'initiatives network', 'research news', 'close menu modal', 'engage', 'engage', 'discover many connect work industry government notforprofit organisation tackle tough challenge .', 'alumni', 'business partnership collaboration', 'global engagement', 'news media relation', 'public policy', 'schools priority group', 'give', 'close menu modal', ' ',', 'founded heart london london lead multidisciplinary university staff student different country .', ' ', 'faculties', 'governance', 'president provost', 'strategy', 'close menu modal', 'active parent page', 'bartlett faculty built environment', 'study', 'research', 'active parent page', 'school institute', 'people', 'ideas', 'engage', 'news events', ' ', 'open close faculty menu', 'breadcrumb trail', 'bartlett faculty built environment', 'school institute', 'bartlett centre advanced spatial analysis', 'faculty menu', 'current page', ' ', 'cupum', 'partnerships', 'software', 'remembering sinesio alves junior', ' ', 'centre advanced spatial analysis . casa . interdisciplinary research institute focus science city within bartlett faculty built environment .', 'centre advanced spatial analysis . casa . establish lead development science city draw upon method idea modelling data science sense urban environment visualisation computation . advancing urban science research collaboration today research still push boundary create good city everyone lead intellectual agenda work closely government industry

partner make realworld impact . teaching reflect make cuttingedge research tool form data longstanding nonacademic partnership train next generation urban scientist skill idea need impact industry government academia third sector . casa community closely connect strongly interdisciplinary . bring together people around world unique variety background include physicist planner geographer economist data scientist architect mathematician computer scientist unite mission tackle challenge face city society around world . work across multiple scale hyperlocal environment lowpowered sensor satellite remote sensing whole country region . studying casa studying casa bring lifelong value student poise take leadership integration role forefront urban spatial data science . study become part active engage alumnus community access listing network social activity well continue contact outstanding teacher researcher . find connected environments casa working casa employ around permanent academic professional service research staff researcher student alumni work edge interdisciplinary science . research tool model frequently inform system range urban plan museum retail heritage application onwards social network understanding . vacancies department publish jobs page . connect linkedin', 'contact', 'tottenham court road', 'london', ' ', 'centre advanced spatial analysis', 'click email .', 'casa ucl.ac.uk', 'click call .', ' . ', 'footer', 'visit', 'bloomsbury theatre studio', 'library museums collections', 'maps', 'shop', 'contact', 'students', 'accommodation', 'current students', 'moodle', 'students union', 'staff', 'inside', 'staff intranet', 'work', 'human resources', 'social medium menu', 'link soundcloud', 'link flickr', 'link tiktok', 'link youtube', 'link instagram', 'link facebook', 'link twitter', 'university college london gower street london', ' . ', . ', 'essential', 'disclaimer', 'freedom information', 'accessibility', 'cookies', 'privacy', 'slavery statement', ' ']

```
help(normalise_document)
```

```
Help on function normalise_document in module textual:

normalise_document(doc: str, html_stripping=True, contraction_expansion=True, accented_c
    Apply all of the functions above to a document using their
    default values so as to demonstrate the NLP process.

    doc: a document to clean.
```

4.2.1 Questions

Let's assume that you want to analyse web page content...

- Based on the above output, what stopwords do you think are missing?
- Based on the above output, what should be removed but isn't?
- Based on the above output, how do you think a computer can work with this text?

> **Stop!**
>
> Beyond this point, we are moving into Natural Language Processing. If you are already struggling with regular expressions, I would recommend *stopping here*. You can come back to revisit the NLP components and creation of word clouds later.

5 Revenons à Nos Moutons

Now that you've seen how the steps are applied to a 'random' HTML document, let's get back to the problem at hand (revenons à nos moutons == let's get back to our sheep).

5.1 Process the Selected Listings

> **Difficulty level: Low, but you'll need to be patient!**

Notice the use of `%%time` here – this will tell you how long each block of code takes to complete. It's a really useful technique for reminding *yourself* and others of how long something might take to run. I find that with NLP this is particularly important since you have to do a *lot* of processing on each document in order to normalise it.

> **Tip**
>
> Notice how we can change the default parameters for `normalise_document` even when using `apply`, but that the syntax is different. So whereas we'd use `normalise_document(doc, remove_digits=True)` if calling the function directly, here it's `.apply(normalise_document, remove_digits=True)`!

Question

```
%%time
# I get about 1 minute on a M2 Mac
lux['description_norm'] = lux.???.apply(???, remove_digits=True)
```

```
%%time
# I get about 1 minute on a M2 Mac
aff['description_norm'] = aff.???.apply(???, remove_digits=True)
```

```
%%time
# I get about 2 seconds on a M2 Mac
bluesp['description_norm'] = bluesp.???.apply(???, remove_digits=True)
```

Answer

```
%%time
# I get about 1 minute on a M2 Mac
lux['description_norm'] = lux.description.apply(normalise_document, remove_digits=Tr
```

```
/Users/jreades/Documents/git/fsds/practicals/textual/__init__.py:606: MarkupResemblesLo

The input looks more like a filename than markup. You may want to open this file and pass the

CPU times: user 43.6 s, sys: 479 ms, total: 44.1 s
Wall time: 44.1 s
```

```
%%time
# I get about 1 minute on a M2 Mac
aff['description_norm'] = aff.description.apply(normalise_document, remove_digits=Tr
```

```
/Users/jreades/Documents/git/fsds/practicals/textual/__init__.py:606: MarkupResemblesLo

The input looks more like a filename than markup. You may want to open this file and pass the

CPU times: user 36.8 s, sys: 431 ms, total: 37.3 s
Wall time: 37.3 s
```

```
%%time
# I get about 1 seconds on a M2 Mac
bluesp['description_norm'] = bluesp.description.apply(normalise_document, remove_dig
```

```
CPU times: user 1.65 s, sys: 20.1 ms, total: 1.67 s
Wall time: 1.68 s
```

5.2 Select and Tokenise

Difficulty level: Low, except for the double list-comprehension.

5.2.1 Select and Extract Corpus

See useful tutorial here. Although we shouldn't have any empty descriptions, by the time we've finished normalising the textual data we may have *created* some empty values and we need to ensure that we don't accidentally pass a NaN to the vectorisers and frequency distribution functions.

```
srcdf = bluesp
```

> **Coding Tip**
>
> Notice how you only need to change the value of the variable here to try any of the different selections we did above? This is a simple kind of parameterisation somewhere between a function and hard-coding everything.

```
corpus = srcdf.description_norm.fillna(' ').values
print(corpus[0:3])
```

```
['house garden close thames river . walk  private road river nearby .  district line undergr
 'space appartment upper floor modernised secure building near canary wharf  fantastic view
 'newly renovate  totally equipped furnished modern apartment heart london . easily accessi
```

5.2.2 Tokenise

There are different forms of tokenisation and different algorithms will expect differing inputs. Here are two:

```
sentences = [nltk.sent_tokenize(text) for text in corpus]
words     = [[nltk.tokenize.word_tokenize(sentence)
                for sentence in nltk.sent_tokenize(text)]
                for text in corpus]
```

Notice how this has turned every sentence into an array and each document into an array of arrays:

```
print(f"Sentences 0: {sentences[0]}")
print()
print(f"Words 0: {words[0]}")
```

```
Sentences 0: ['house garden close thames river .', 'walk  private road river nearby .', 'dis

Words 0: [['house', 'garden', 'close', 'thames', 'river', '.'], ['walk', 'private', 'road'
```

5.3 Frequencies and Ngrams

> **Difficulty level: Moderate.**

One new thing you'll see here is the `ngram`: ngrams are 'simply' pairs, or triplets, or quadruplets of words. You may come across the terms unigram (`ngram(1,1)`), bigram (`ngram(2,2)`), trigram (`ngram(3,3)`)... typically, you will rarely find anything beyond trigrams, and these present real issues for text2vec algorithms because the embedding for `geographical`, `information`, and `systems` is *not* the same as for `geographical information systetms`.

5.3.1 Build Frequency Distribution

Build counts for ngram range 1..3:

```
fcounts = dict()

# Here we replace all full-stops... can you think why we might do this?
data = nltk.tokenize.word_tokenize(' '.join([text.replace('.','') for text in corpus

for size in 1, 2, 3:
    fdist = FreqDist(ngrams(data, size))
    print(fdist)
    # If you only need one note this: https://stackoverflow.com/a/52193485/4041902
    fcounts[size] = pd.DataFrame.from_dict({f'Ngram Size {size}': fdist})
```

```
<FreqDist with 2692 samples and 26245 outcomes>
<FreqDist with 14173 samples and 26244 outcomes>
<FreqDist with 19540 samples and 26243 outcomes>
```

5.3.2 Output Top-n Ngrams

And output the most common ones for each ngram range:

```
for dfs in fcounts.values():
    print(dfs.sort_values(by=dfs.columns.values[0], ascending=False).head(10))
    print()
```

```
           Ngram Size 1
walk                594
room                472
london              469
river               418
bedroom             412
space               391
minute              382
apartment           373
station             307
flat                303

                 Ngram Size 2
minute  walk              252
river   thames            138
        view              138
living  room              134
canary  wharf             112
guest   access            110
central london            107
fully   equip              76
equip   kitchen            71
```

```
thames  river              65

                        Ngram Size 3
fully   equip   kitchen            68
walk    river   thames             37
close   river   thames             35
walk    thames  river              27
minute  walk    river              23
open    plan    kitchen            20
thames  river   view               20
sofa    living  room               19
within  walk    distance           19
open    plan    live               18
```

5.3.3 Questions

- Can you think why we don't care about punctuation for frequency distributions and n-grams?
- Do you understand what n-grams *are*?

5.4 Count Vectoriser

Difficulty level: Low, but the output needs some thought!

This is a big foray into sklearn (sci-kit learn) which is the main machine learning and clustering module for Python. For processing text we use *vectorisers* to convert terms to a vector representation. We're doing this on the smallest of the derived data sets because these processes can take a while to run and generate *huge* matrices (remember: one row and one column for each term!).

5.4.1 Fit the Vectoriser

```
cvectorizer = CountVectorizer(ngram_range=(1,3))
cvectorizer.fit(corpus)
```

```
CountVectorizer(ngram_range=(1, 3))
```

5.4.2 Brief Demonstration

Find the number associated with a word in the vocabulary and how many times it occurs in the original corpus:

```
term = 'stratford'
pd.options.display.max_colwidth=750
# Find the vocabulary mapping for the term
```

```
print(f"Vocabulary mapping for {term} is {cvectorizer.vocabulary_[term]}")
# How many times is it in the data
print(f"Found {srcdf.description_norm.str.contains(term).sum():,} rows containing {t
# Print the descriptions containing the term
for x in srcdf[srcdf.description_norm.str.contains(term)].description_norm:
    as_markdown('Stratford',x)
```

```
Vocabulary mapping for stratford is 29373
Found 10 rows containing stratford
```

Stratford

house garden close thames river . walk private road river nearby . district line underground . walk . direct access central london near gardens . kids playground walk distance along thames path . space residential neighborhood english corporate expat family . house culdesac private road river thames . river foot away . walking distance subway . central london underground district line . gardens stop walk zone . addition overground stratford also stop gardens underground station . gardens stop walk . overland railway station bridge . walk . take waterloo railway station minute . bicycle follow towpath hammersmith bridge continue putney bridge . lastly several stree

Stratford

please read things note comfortable clean bright brand flat east london minute central london tube quite central great transport link major london attraction minute walk river park undergroundtubedlr station supermarket docklands stratford olympic stadium westfield shopping centre . enjoy brick lane indian restaurant spitalfields market colombian flower market historical whitechapel . space please read things note nice clean fresh bright airy . space perfect professional single person couple . make feel like home choice anything like wake relaxing cooking . guest access please read things note entire flat . please treat home away . please treat .

Stratford

comfortable fairly flat east london travel zone minute central london quite central great transport link major london attraction minute walk river park undergroundtube station minute supermarket minute docklands stratford olympic stadium westfield shopping centre . enjoy brick lane indian restaurant spitalfields market colombian flower market historical whitechapel . space spacious comfortable tidy clean airy relaxing . live flat sleep open plan lounge balcony . guest access bathroom share . welcome microwave ready meal toaster make drink till fridge store food . please make sure clean clear immediately . dining table . stay present morning evening weekend . also

Stratford

entire appartment double bedroom large living area.the apartment feature kitchen come free wifi flat screen tv.to make exicting luxurious even free free sauna well . space stunning apartment london docklands bank thames river close thames barrier park canary wharf . apartment floor spacious living room balcony . nearest station pontoon dock walkable distance direct train stratford . mins . bank . excel centre . walk . arena canary wharf mins train mins central london . world heritage sitethames barrier thames barrier park walk appartment . london city airport train station away . fully kitchen bathroom broadband internet underground secure parking onsite . attract

Stratford

luxurious bedroom apartment zone love hidden secret part town minute away everywhere river view slow pace main artery town right doorstep well hidden beauty park waterway . easy . walk tube route center town well stratford olympic park canary wharf much much right doorstep space welcome home place love bedroom hold personal belonging bedroom give guest idea size . bedroom large double accommodate comfortably . sofa chair accommodate guest extra extra charge . welcome guest personally wish know . therefore important check time convenient . midnight arrival . time ehich discuss good

Stratford

place close mile tube station brick lane shoreditch queen mary university london stratford westfield minute tube central london . love place newly renovate flat amazing canal view guest bedroom clean friendly environment . place good couple solo adventurer business traveller .

Stratford

locate high street give amazing water view stadium sight amazing architectural structure walk pudding mill lane walk abba walk stratford westfield walk stratfordstratford international station mins walk mins train ride central london

Stratford

modern spacious bedroom suite apartment close river thames wimbledon . situate wandsworth district london england lawn tennis club centre court .km clapham junction . stratford bridge chelsea . city view free wifi throughout property . apartment feature bedroom kitchen fridge oven wash machine flat screen seating area bathroom shower . eventim .km away .

Stratford

> perfect group trip . modern spacious suite apartment close river thames wimbledon . situated wandsworth district london england lawn tennis club centre court .km clapham junction . stratford bridge chelsea . city view free wifi throughout property . apartment feature bedroom kitchen wfridge oven wash machine flat screen seating area bathroom wshower . eventim .km away .

Stratford

> flat locate zone east london near canary wharf . nice quiet residential area canal . flat amazing canal view balcony . enjoy morning coffee swan goose everyday . huge park opposite flat picnic . canary wharf shop mall . mins bank stratford westfield . mins central oxford circus tube . locate convenient transportation link .

5.4.3 Transform the Corpus

You can only *tranform* the entire corpus *after* the vectoriser has been fitted. There is an option to `fit_transform` in one go, but I wanted to demonstrate a few things here and some vectorisers are don't support the one-shot fit-and-transform approach. **Note the type of the transformed corpus**:

```
cvtcorpus = cvectorizer.transform(corpus)
cvtcorpus # cvtcorpus for count-vectorised transformed corpus
```

```
<408x35278 sparse matrix of type '<class 'numpy.int64'>'
    with 71420 stored elements in Compressed Sparse Row format>
```

5.4.4 Single Document

Here is the **first** document from the corpus:

```
doc_df = pd.DataFrame(cvtcorpus[0].T.todense(),
                      index=cvectorizer.get_feature_names_out(), columns=["Counts"]
                     ).sort_values('Counts', ascending=False)
doc_df.head(10)
```

	Counts
walk	6
gardens	4
river	4
bridge	3
stop	3
thames	3
station	3

	Counts
underground	3
railway	2
central london	2

5.4.5 Transformed Corpus

```
cvdf = pd.DataFrame(data=cvtcorpus.toarray(),
                    columns=cvectorizer.get_feature_names_out())
print(f"Raw count vectorised data frame has {cvdf.shape[0]:,} rows and {cvdf.shape[1
cvdf.iloc[0:5,0:10]
```

```
Raw count vectorised data frame has 408 rows and 35,278 columns.
```

	aaathe	aaathe apartment	aaathe apartment quiet	aand	aand comfy	aand comfy sofa
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

5.4.6 Filter Low-Frequency Words

These are likely to be artefacts of text-cleaning or human input error. As well, if we're trying to look across an entire corpus then we might not want to retain words that only appear in a couple of documents.

Let's start by getting the *column* sums:

```
sums = cvdf.sum(axis=0)
print(f"There are {len(sums):,} terms in the data set.")
sums.head()
```

```
There are 35,278 terms in the data set.
```

```
aaathe                  1
aaathe apartment        1
aaathe apartment quiet  1
aand                    1
aand comfy              1
dtype: int64
```

Remove columns (i.e. terms) appearing in less than 1% of documents. You can do this by thinking about what the shape of the data frame means (rows and/or columns) and how you'd get 1% of that!

Question

```
filter_terms = sums >= cvdf.shape[0] * ???
```

Answer

```
filter_terms = sums >= cvdf.shape[0] * 0.01
```

Now see how we can use this to strip out the columns corresponding to low-frequency terms:

```
fcvdf = cvdf.drop(columns=cvdf.columns[~filter_terms].values)
print(f"Filtered count vectorised data frame has {fcvdf.shape[0]:,} rows and {fcvdf.
fcvdf.iloc[0:5,0:10]
```

```
Filtered count vectorised data frame has 408 rows and 2,043 columns.
```

	able	access	access access	access bathroom	access central	access central london	ac
0	0	1	0	0	1	1	0
1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0
4	0	1	0	0	0	0	0

```
fcvdf.sum(axis=0)
```

```
able                          8
access                      242
access access                 7
access bathroom               5
access central                8
                            ...
zone comfortable cosy         7
zone near                     5
zone near underground         5
zone recently                10
zone recently refurbish      10
Length: 2043, dtype: int64
```

We're going to pick this up again in Task 7.

5.4.7 Questions

- Can you explain what `doc_df` contains?
- What does `cvdf` contain? Explain the rows and columns.
- What is the function of `filter_terms`?

5.5 TF/IDF Vectoriser

> ⚠ Difficulty level: Moderate
>
> But only if you want to understand how `max_df` and `min_df` work!

5.5.1 Fit and Transform

```
tfvectorizer = TfidfVectorizer(use_idf=True, ngram_range=(1,3),
                               max_df=0.75, min_df=0.01) # <-- these matter!
tftcorpus    = tfvectorizer.fit_transform(corpus) # TF-transformed corpus
```

5.5.2 Single Document

```
doc_df = pd.DataFrame(tftcorpus[0].T.todense(), index=tfvectorizer.get_feature_names
doc_df.sort_values('Weights', ascending=False).head(10)
```

	Weights
gardens	0.414885
stop	0.241659
district line	0.239192
railway	0.232131
underground	0.201738
district	0.197221
bridge	0.191983
walk	0.189485
road	0.151163
distance	0.142999

5.5.3 Transformed Corpus

```
tfidf = pd.DataFrame(data=tftcorpus.toarray(),
                     columns=tfvectorizer.get_feature_names_out())
print(f"TF/IDF data frame has {tfidf.shape[0]:,} rows and {tfidf.shape[1]:,} columns
tfidf.head()
```

```
TF/IDF data frame has 408 rows and 1,911 columns.
```

	able	access	access access	access bathroom	access central	access central london
0	0.0	0.043972	0.0	0.0	0.11031	0.11031
1	0.0	0.000000	0.0	0.0	0.00000	0.00000
2	0.0	0.000000	0.0	0.0	0.00000	0.00000

	able	access	access access	access bathroom	access central	access central london
3	0.0	0.000000	0.0	0.0	0.00000	0.00000
4	0.0	0.044127	0.0	0.0	0.00000	0.00000

5.5.4 Questions

- What does the TF/IDF score *represent*?
- What is the role of `max_df` and `min_df`?

6 Word Clouds

6.1 For Counts

Difficulty level: Easy!

```
fcvdf.sum().sort_values(ascending=False)
```

```
walk                        595
room                        472
london                      471
river                       418
bedroom                     412
                           ...
chic                          5
term                          5
choice                        5
teddington                    5
london aquarium minute        5
Length: 2043, dtype: int64
```

```
ff = 'RobotoMono-VariableFont_wght.ttf'
dp = '/home/jovyan/fonts/'
tp = os.path.join(os.path.expanduser('~'),'Library','Fonts')
if os.path.exists(tp):
    fp = os.path.join(tp,ff)
else:
    fp = os.path.join(dp,ff)
```

```
f,ax = plt.subplots(1,1,figsize=(8, 8))
plt.gcf().set_dpi(150)
Cloud = WordCloud(
    background_color="white",
    max_words=75,
    font_path=fp
```

```
).generate_from_frequencies(fcvdf.sum())
ax.imshow(Cloud)
ax.axis("off");
#plt.savefig("Wordcloud 1.png")
```

6.2 For TF/IDF Weighting

> Difficulty level: Low, but you'll need to be patient!

```
tfidf.sum().sort_values(ascending=False)
```

```
walk                       23.037285
room                       19.135852
london                     18.744519
minute                     18.650909
apartment                  18.082855
                               ...
station apartment one       0.401426
station also close          0.401426
apartment one benefit       0.401426
apartment one               0.401426
also close station          0.401426
Length: 1911, dtype: float64
```

```
f,ax = plt.subplots(1,1,figsize=(8, 8))
plt.gcf().set_dpi(150)
Cloud = WordCloud(
    background_color="white",
    max_words=100,
    font_path=fp
```

```
).generate_from_frequencies(tfidf.sum())
ax.imshow(Cloud)
ax.axis("off");
#plt.savefig("Wordcloud 2.png")
```

6.2.1 Questions

- What does the `sum` represent for the count vectoriser?
- What does the `sum` represent for the TF/IDF vectoriser?

7 Latent Dirchlet Allocation

> **Tip**
>
> I would give this a *low* priority. It's a commonly-used method, but on small data sets it really isn't much use and I've found its answers to be... unclear... even on large data sets.

Adapted from this post on doing LDA using sklearn. Most other examples use the `gensim` library.

```
# Notice change to ngram range
# (try 1,1 and 1,2 for other options)
vectorizer = CountVectorizer(ngram_range=(1,2))
```

7.1 Calculate Topics

```
vectorizer.fit(corpus)
tcorpus = vectorizer.transform(corpus) # tcorpus for transformed corpus

LDA = LatentDirichletAllocation(n_components=3, random_state=42) # Might want to exp
LDA.fit(tcorpus)
```

```
LatentDirichletAllocation(n_components=3, random_state=42)
```

```
first_topic = LDA.components_[0]
top_words = first_topic.argsort()[-25:]

for i in top_words:
    print(vectorizer.get_feature_names_out()[i])
```

```
river thames
modern
area
wharf
flat
access
bathroom
guest
house
minute
private
kitchen
large
thames
living
view
station
floor
bedroom
apartment
walk
london
river
room
space
```

```
for i,topic in enumerate(LDA.components_):
    as_markdown(f'Top 10 words for topic #{i}', ', '.join([vectorizer.get_feature_na
```

Top 10 words for topic #0

river thames, modern, area, wharf, flat, access, bathroom, guest, house, minute, private, kitchen, large, thames, living, view, station, floor, bedroom, apartment, walk, london, river, room, space

Top 10 words for topic #1

park, fully, modern, private, living, guest, view, double, close, access, area, flat, thames, bathroom, station, apartment, kitchen, space, minute walk, room, london, river, bedroom, minute, walk

Top 10 words for topic #2

living, river view, canary wharf, canary, central, wharf, stay, close, bathroom, guest, kitchen, double, minute, thames, access, station, flat, space, view, bedroom, apartment, river, walk, london, room

7.2 Maximum Likelihood Topic

```
topic_values = LDA.transform(tcorpus)
topic_values.shape
```

```
(408, 3)
```

```
pd.options.display.max_colwidth=20
srcdf['Topic'] = topic_values.argmax(axis=1)
srcdf.head()
```

	geometry	listing_url	name	description	ame
19	POINT (519474.79...	https://www.airb...	Townhouse in Ric...	3 Bed House with...	["Bat
71	POINT (537104.16...	https://www.airb...	Rental unit in L...	The space...	["Hea
713	POINT (530945.88...	https://www.airb...	Rental unit in G...	Newly renovated,...	["Bat
928	POINT (539808.31...	https://www.airb...	Rental unit in L...	Brand new rivers...	["Wa
1397	POINT (538098.29...	https://www.airb...	Rental unit in L...	PLEASE READ OTHE...	["Hea

```
pd.options.display.max_colwidth=75
srcdf[srcdf.Topic==1].description_norm.head(10)
```

```
19     house garden close thames river . walk  private road river nearby .  di...
928    brand riverside apartment greenwich peninsula . perfect explore london ...
1424   mint walk thames river  mint tower bridge  mint greenwich bus mint walk...
1644   bright  spacious chill bedroom cosy apartment level friendly quiet bloc...
3887   fabulous bathshower flat modern development east putney  london literal...
3966   large double room river view available large spacious townhouse hampton...
```

```
4664   private quiet bright  large double room ensuite bathroom . location stu...
4890   bright  airy bedroom  ground floor apartment quiet street . modernly fu...
5815   perfect claphambattersea modern decor . brand kitchenbathroom minute wa...
6699   room front house overlook balcony  road river . bright sunny house view...
Name: description_norm, dtype: object
```

```
vectorizer = CountVectorizer(ngram_range=(1,1), stop_words='english', analyzer='word
topic_corpus = vectorizer.fit_transform(srcdf[srcdf.Topic==1].description.values) #
```

```
topicdf = pd.DataFrame(data=topic_corpus.toarray(),
                       columns=vectorizer.get_feature_names_out())
```

```
f,ax = plt.subplots(1,1,figsize=(8,8))
plt.gcf().set_dpi(150)
Cloud = WordCloud(
            background_color="white",
            max_words=75).generate_from_frequencies(topicdf.sum())
ax.imshow(Cloud)
ax.axis("off");
# plt.savefig('Wordcloud 3.png')
```

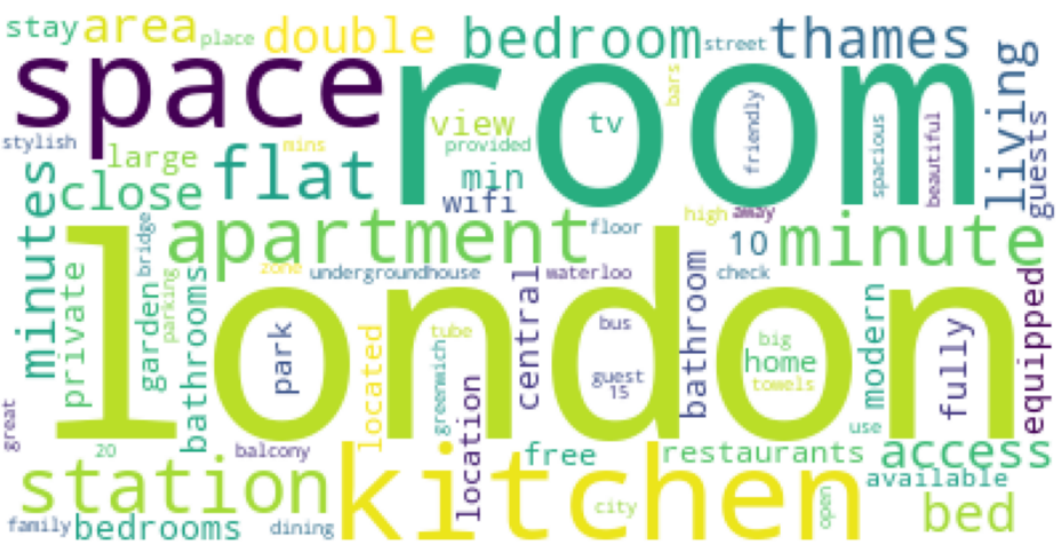

8 Word2Vec

> **Tip**
>
> This algorithm works almost like magic. You should play with the configuration parameters and see how it changes your results.

8.1 Configure

```
from gensim.models.word2vec import Word2Vec
```

```
dims = 100
print(f"You've chosen {dims} dimensions.")

window = 4
print(f"You've chosen a window of size {window}.")

min_v_freq  = 0.01 # Don't keep words appearing less than 1% frequency
min_v_count = math.ceil(min_v_freq * srcdf.shape[0])
print(f"With a minimum frequency of {min_v_freq} and {srcdf.shape[0]:,} documents, m
```

```
You've chosen 100 dimensions.
You've chosen a window of size 4.
With a minimum frequency of 0.01 and 408 documents, minimum vocab frequency is 5.
```

8.2 Train

```
%%time

corpus      = srcdf.description_norm.fillna(' ').values
#corpus_sent = [nltk.sent_tokenize(text) for text in corpus] # <-- with more formal
corpus_sent = [d.replace('.',' ').split(' ') for d in corpus] # <-- deals better wit
model       = Word2Vec(sentences=corpus_sent, vector_size=dims, window=window, epoch
                min_count=min_v_count, seed=42, workers=1)

#model.save(f"word2vec-d{dims}-w{window}.model") # <-- You can then Word2Vec.load(..
```

8.3 Explore Similarities

This next bit of code only runs if you have calculated the frequencies above in the Frequencies and Ngrams section.

```
pd.set_option('display.max_colwidth',150)

df = fcounts[1] # <-- copy out only the unigrams as we haven't trained anything else

n     = 14 # number of words
topn  = 7  # number of most similar words

selected_words = df[df['Ngram Size 1'] > 5].reset_index().level_0.sample(n, random_s

words = []
v1    = []
v2    = []
```

```
v3    = []
sims  = []

for w in selected_words:
    try:
        vector = model.wv[w]  # get numpy vector of a word
        #print(f"Word vector for '{w}' starts: {vector[:5]}...")

        sim = model.wv.most_similar(w, topn=topn)
        #print(f"Similar words to '{w}' include: {sim}.")

        words.append(w)
        v1.append(vector[0])
        v2.append(vector[1])
        v3.append(vector[2])
        sims.append(", ".join([x[0] for x in sim]))
    except KeyError:
        print(f"Didn't find {w} in model. Can happen with low-frequency terms.")

vecs = pd.DataFrame({
    'Term':words,
    'V1':v1,
    'V2':v2,
    'V3':v3,
    f'Top {topn} Similar':sims
})

vecs
```

```
#print(model.wv.index_to_key) # <-- the full vocabulary that has been trained
```

8.4 Apply

We're going to make *use* of this further next week...

8.4.1 Questions

- What happens when *dims* is very small (e.g. 25) or very large (e.g. 300)?
- What happens when *window* is very small (e.g. 2) or very large (e.g. 8)?

9 Processing the Full File

> **Caution**
>
> This code can take *some time* (**> 5 minutes on a M2 Mac**) to run, so **don't run this** until you've understood what we did before!

You will get a warning about `"." looks like a filename, not markup` — this looks a little scary, but is basically suggesting that we have a description that consists only of a '.' or that looks like some kind of URL (which the parser thinks means you're trying to pass it something to download).

```
%%time
# This can take up to 8 minutes on a M2 Mac
gdf['description_norm'] = ''
gdf['description_norm'] = gdf.description.apply(normalise_document, remove_digits=Tr
```

```
gdf.to_parquet(os.path.join('data','geo',f'{fn.replace(".","-with-nlp.")}'))
```

> **Tip**
>
> Saving an intermediate file at this point is useful because you've done quite a bit of *expensive* computation. You *could* restart-and-run-all and then go out for the day, but probably easier to just save this output and then, if you need to restart your analysis at some point in the future, just remember to deserialise amenities back into a list format.

9.1 Applications

The above is *still* only the results for the one of the subsets of apartments *alone*. At this point, you would probably want to think about how your results might change if you changed any of the following:

1. Using one of the other data sets that we created, or even the entire data set!
2. Applying the CountVectorizer or TfidfVectorizer *before* selecting out any of our 'sub' data sets.
3. Using the visualisation of information to improve our regex selection process.
4. Reducing, increasing, or constraining (i.e. `ngrams=(2,2)`) the size of the ngrams while bearing in mind the impact on processing time and interpretability.
5. Filtering by type of listing or host instead of keywords found in the description (for instance, what if you applied TF/IDF to the entire data set and then selected out 'Whole Properties' before splitting into those advertised by hosts with only one listing vs. those with multiple listings?).
6. Linking this back to the geography.

Over the next few weeks we'll also consider alternative means of visualising the data!

9.2 Resources

There is a lot more information out there, including a whole book and your standard O'Reilly text.

And some more useful links:

- Pandas String Contains Method
- Using Regular Expressions with Pandas
- Summarising Chapters from Frankenstein using TF/IDF